

L0. Introduction

Motivation

- Machine Learning
- Artificial Intelligence
- Optimization
- Optimum Control
- Planning
- Markov Decision Process
- Influential Diagram
- Decision Tree
- Dynamic Control
- Game Theory
- Search
- Stochastic Programming
- Reinforcement Learning
- Bandit problem
- :

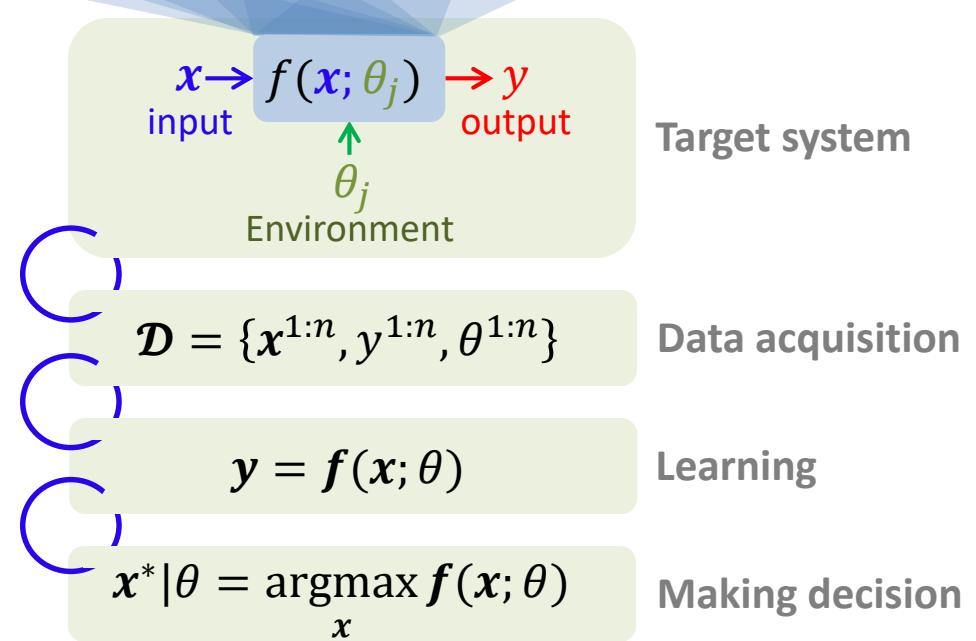
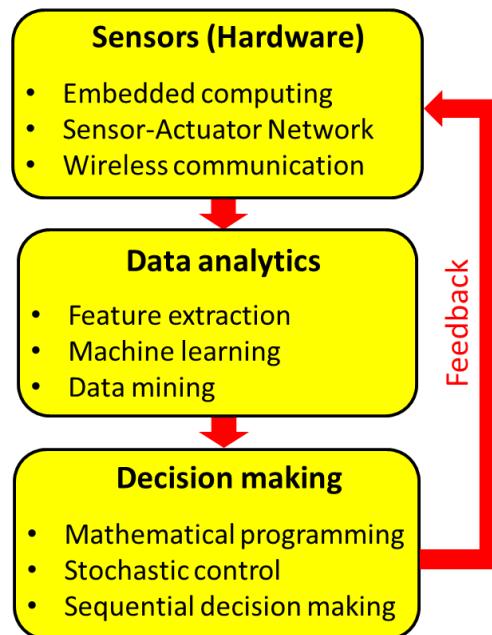
Engineering is all about decision makings



What are **the differences** in these decision-making strategies

What are **the common aspects** in these decision-making strategies?

Data-driven decision making and control in engineering domain



Motivation

engineers' creativity depends on the diversity of tools that he or she has



- Diversify tools for decision making
- Understand the usages of your tools
- Sharpen your tools
- Organize your tools

What type of decision making framework will be used?

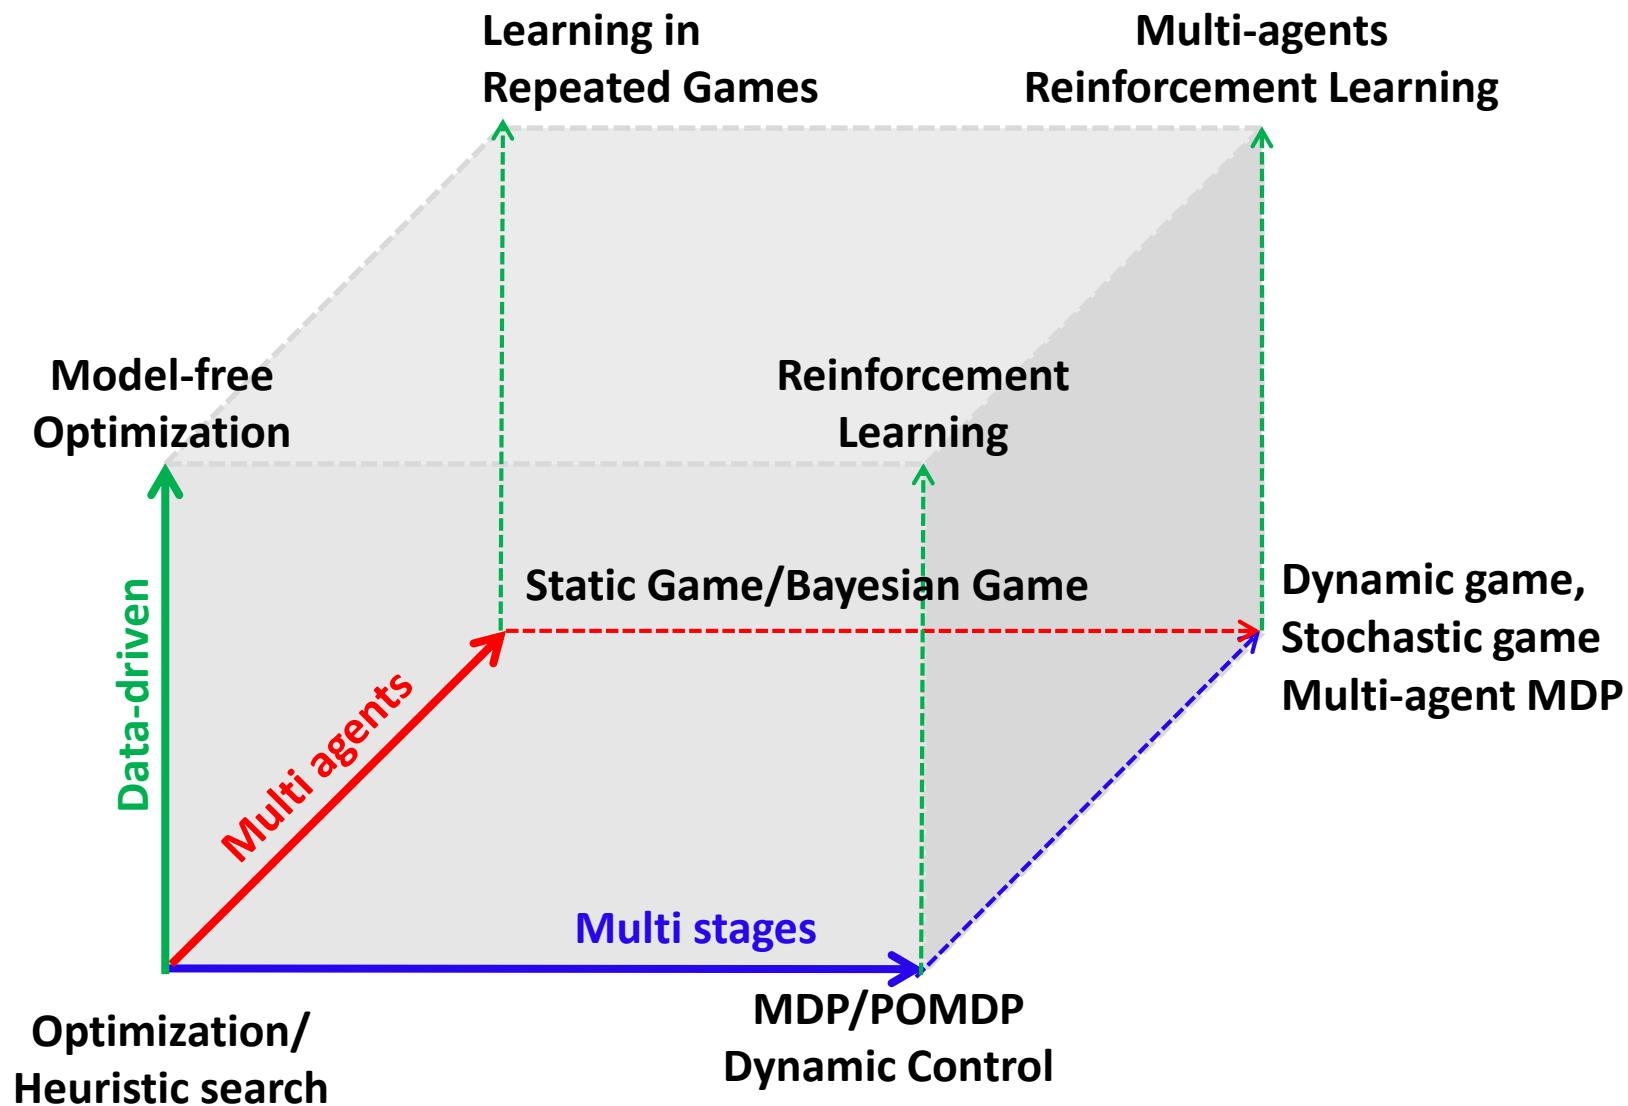
- Single stage or multi stages
- Single decision maker or many decision makers
- Model based or model-free

“Decision makings under uncertainties”

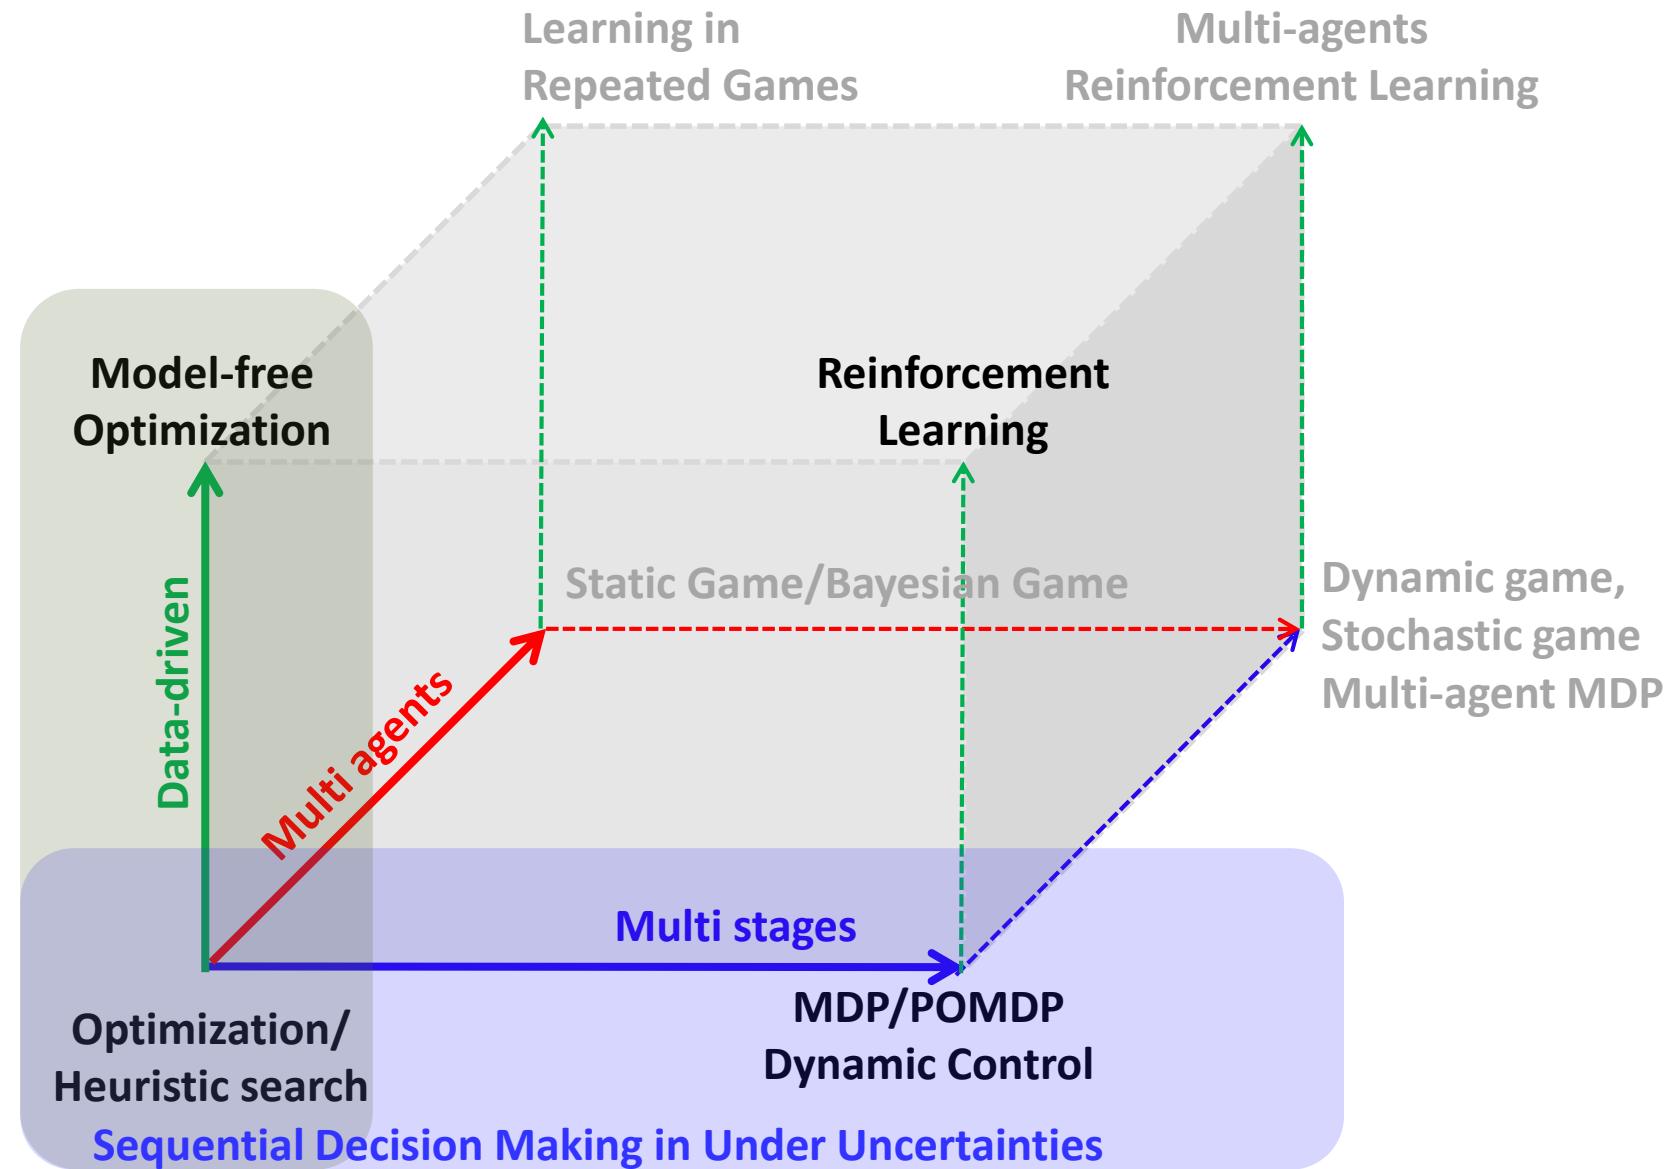
How to model uncertainties?

- **Epistemic Uncertainty** (systemic uncertainty) :
Uncertainty arising through lack of knowledge
 - Model uncertainty
 - State uncertainty
- **Aleatoric uncertainty** (statistical uncertainty):
Uncertainty arising through an underlying stochastic system

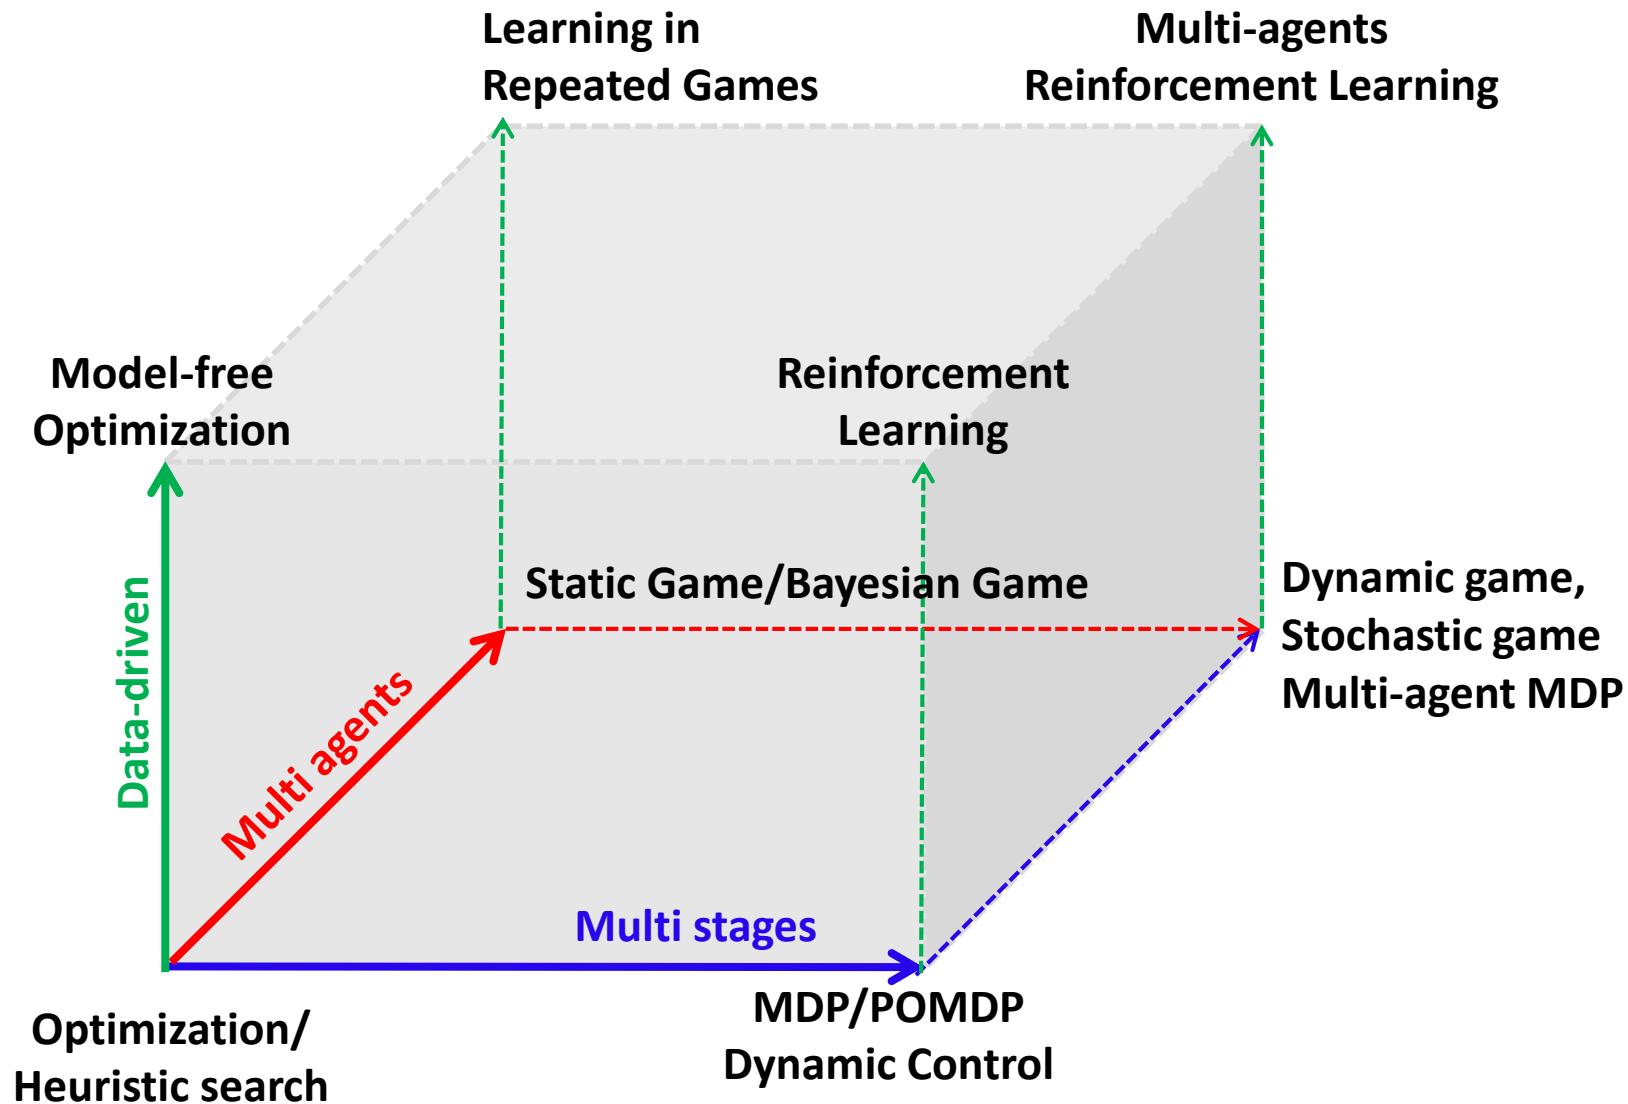
Classified decision making methods



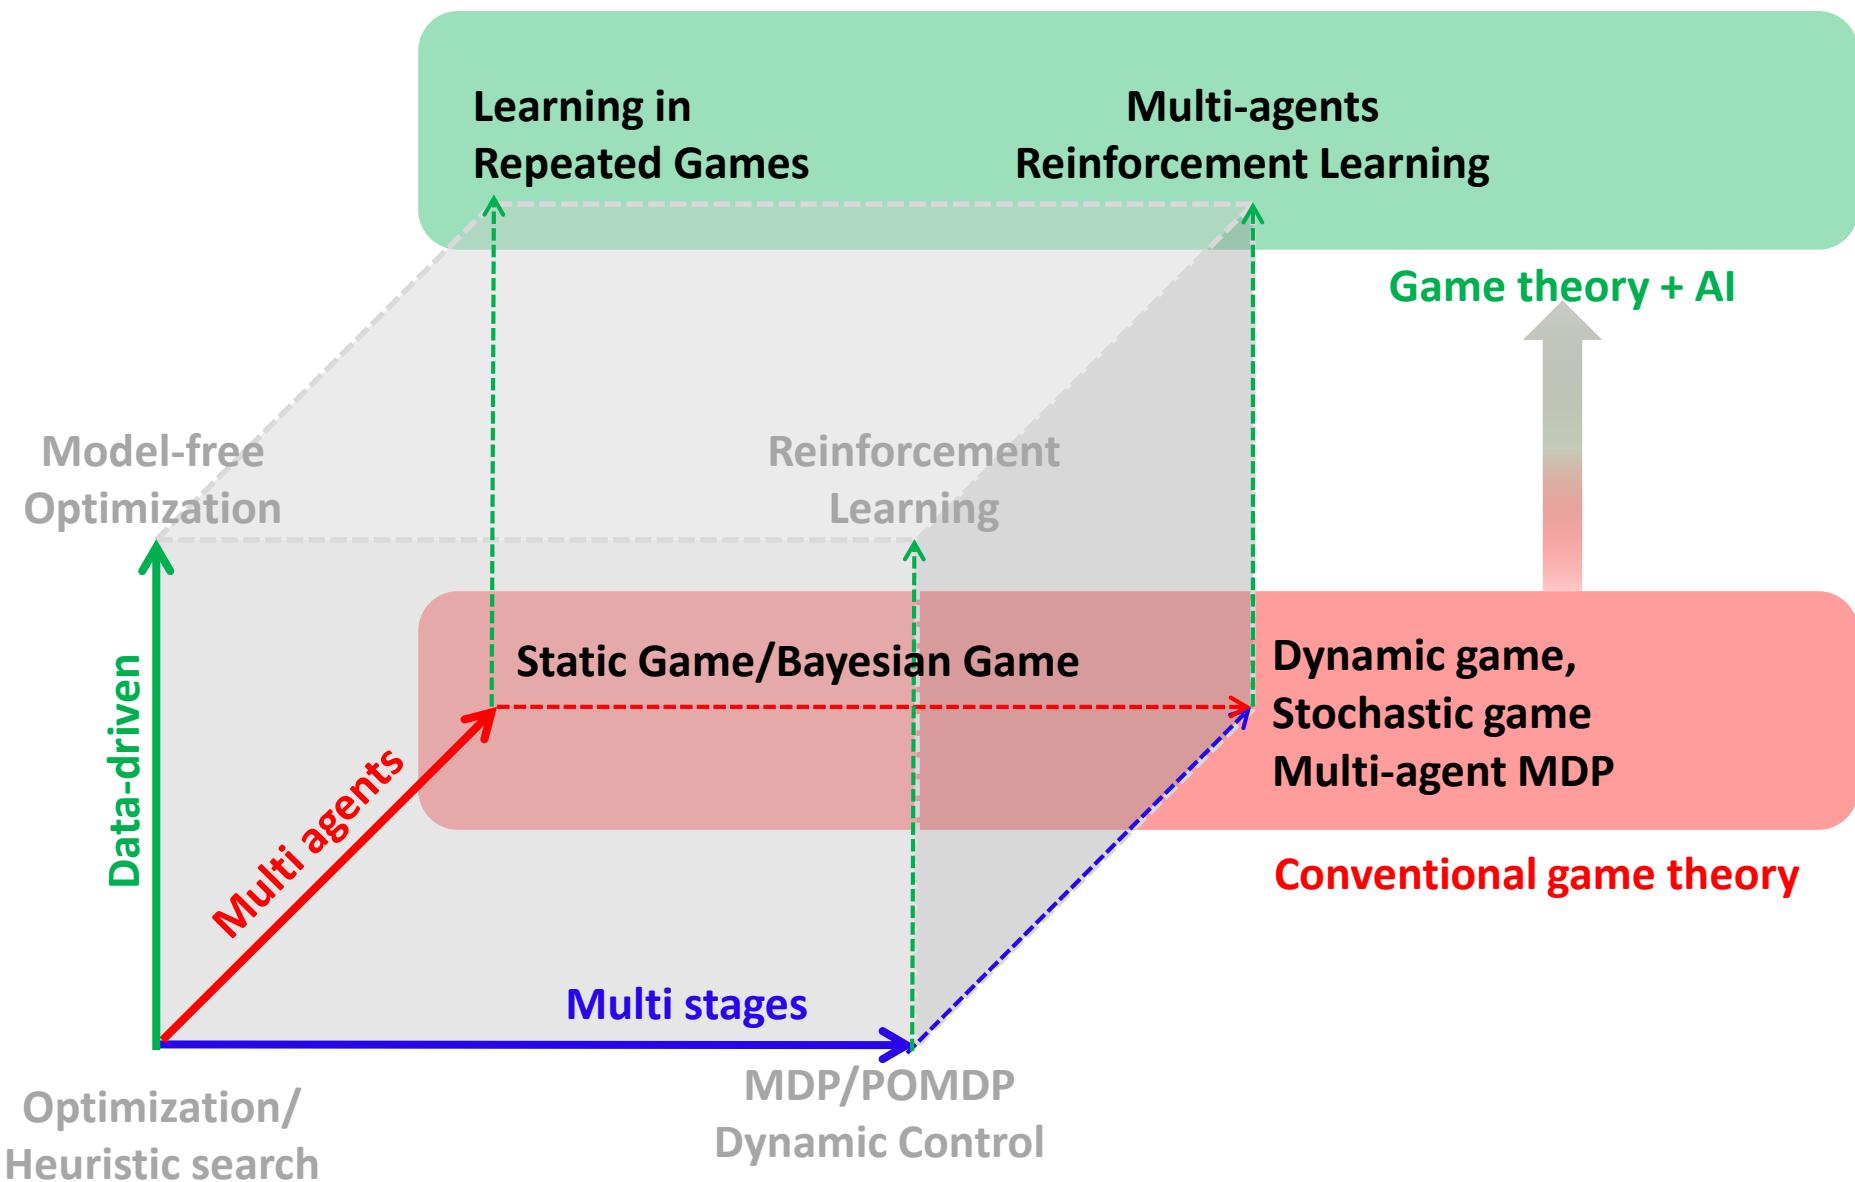
Classified decision making methods



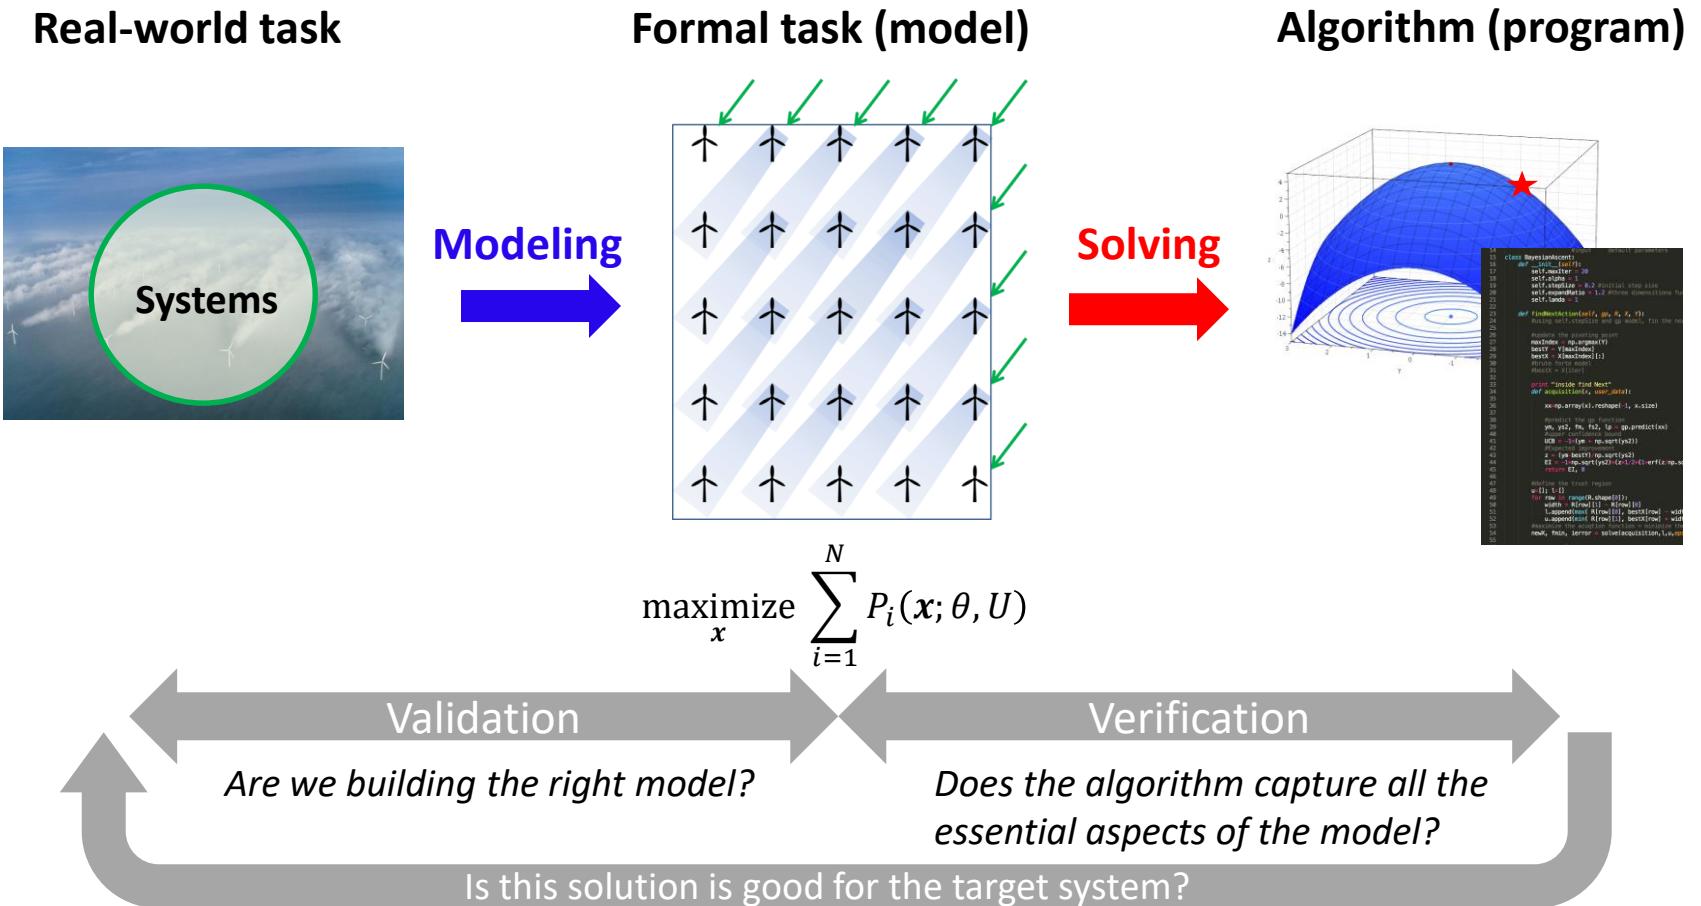
Classified decision making methods



Classified decision making methods



Problem solving



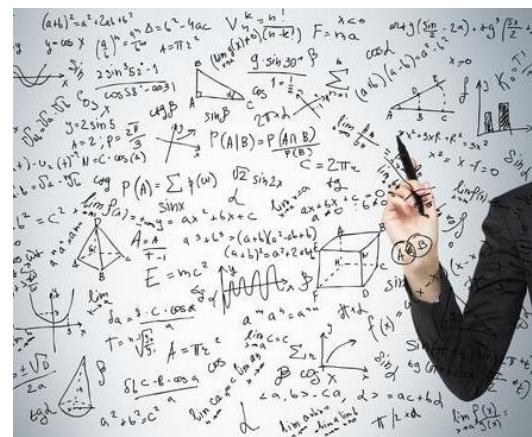
Data can help model more realistically and derive more accurate solution!

Key elements of this course



Data analytics

- Bayesian Statistics
- Machine learning
- Bayesian Network



Modeling

- Optimization
- Markov Decision Process
- Game Theory



Decision Making

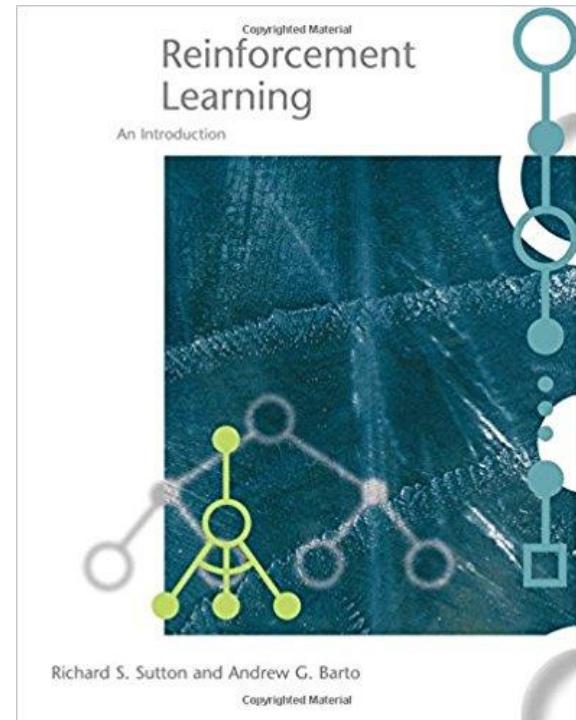
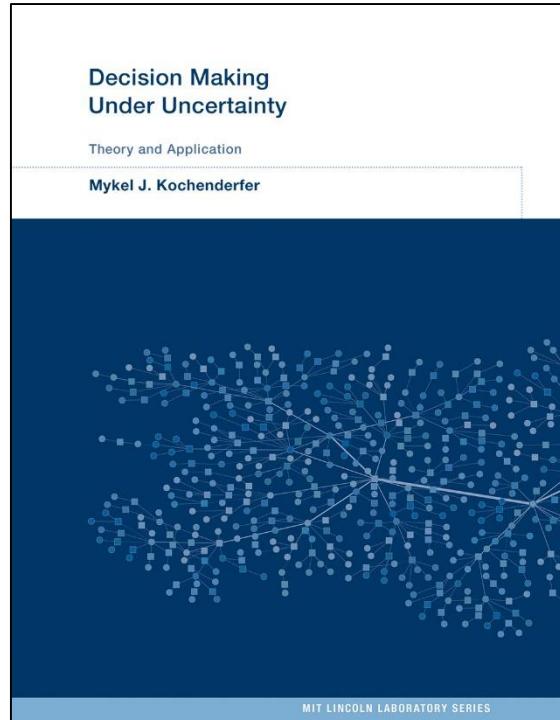
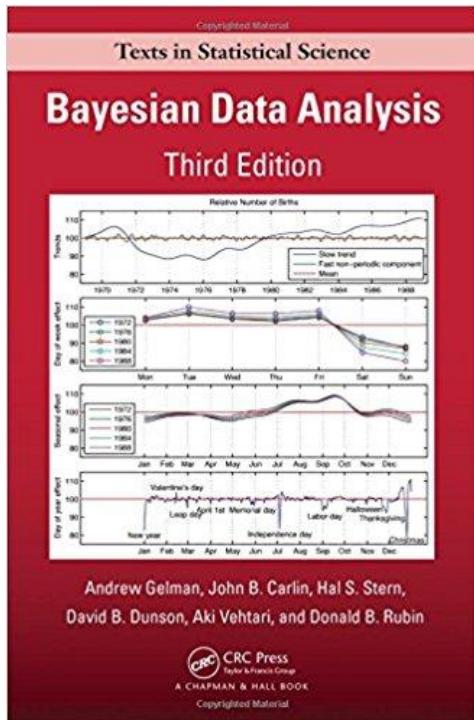
- Mathematical Programming
- Dynamic Programming
- Reinforcement Learning

Course Objectives

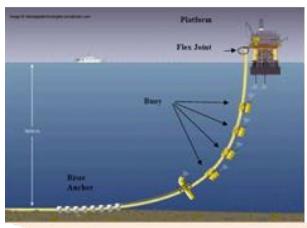
Upon successful completion of the course, you are able to

- *understand* various mathematical models describing decision making problems.
- *formulate* real-world decision making problems in a mathematical form.
- *implement* key algorithms and approaches to solving various decision making problems.
- *Interpret* the results of decision-making problems.

Books (optional)



Engineering Systems are Stochastic-Nonlinear-Dynamical Systems



Engineering Systems are
Stochastic-Nonlinear-Dynamical Systems



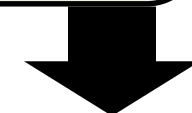
System Identification

- White-box model
- Gray-box model
- Black-box model



System monitoring

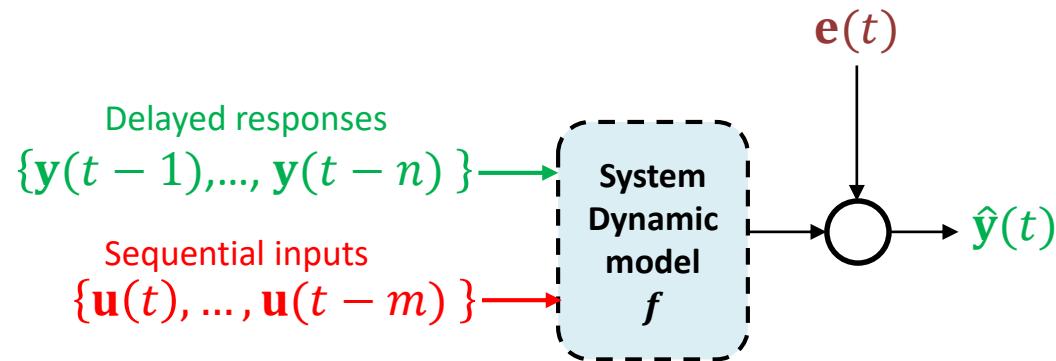
- Condition based monitoring
- Fault detection & prognostic



System Control

- Planning & Scheduling
- Operation & Management

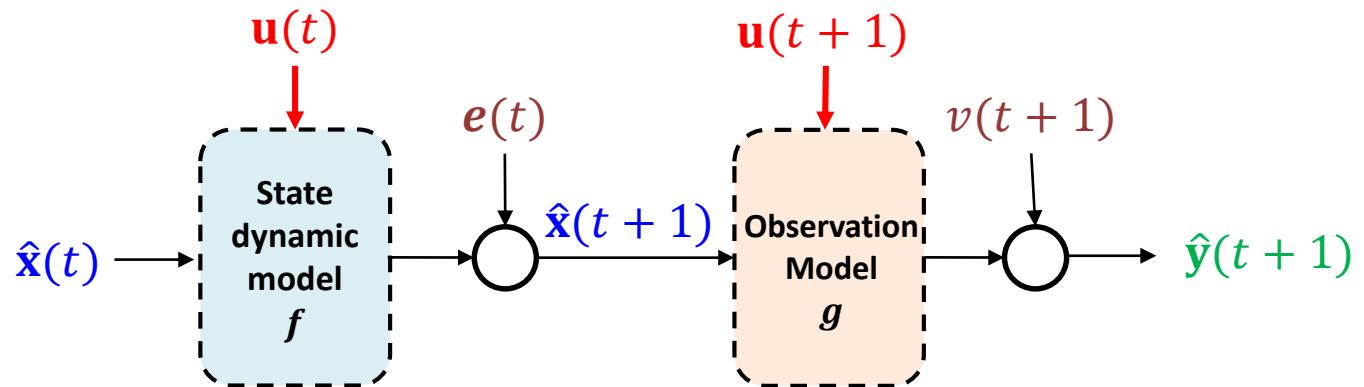
Systems Identification



NARX (nonlinear autoregressive model with exogenous input)

$$\hat{\mathbf{y}}(t) = f(\mathbf{y}(t-1), \dots, \mathbf{y}(t-n), \mathbf{u}(t), \dots, \mathbf{u}(t-m)) + e(t)$$

Systems Identification



State-space model

- State dynamic model

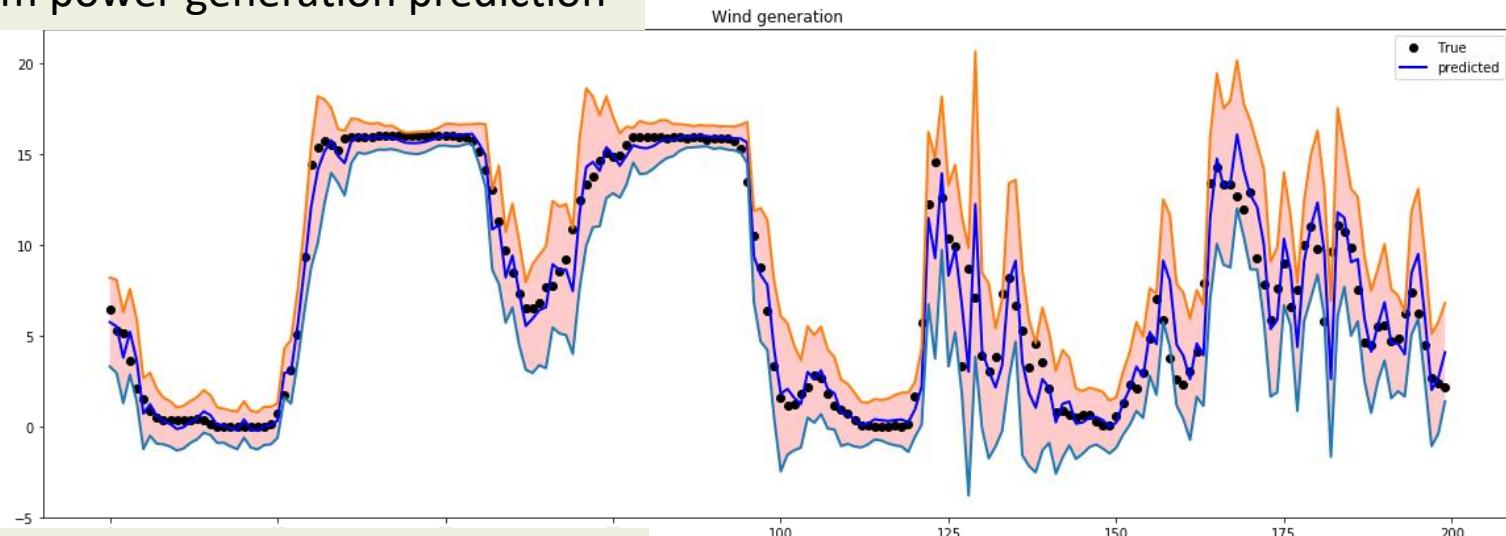
$$\hat{x}(t+1) = f(\mathbf{x}(t), \mathbf{u}(t)) + \mathbf{e}(t)$$

- Observation model

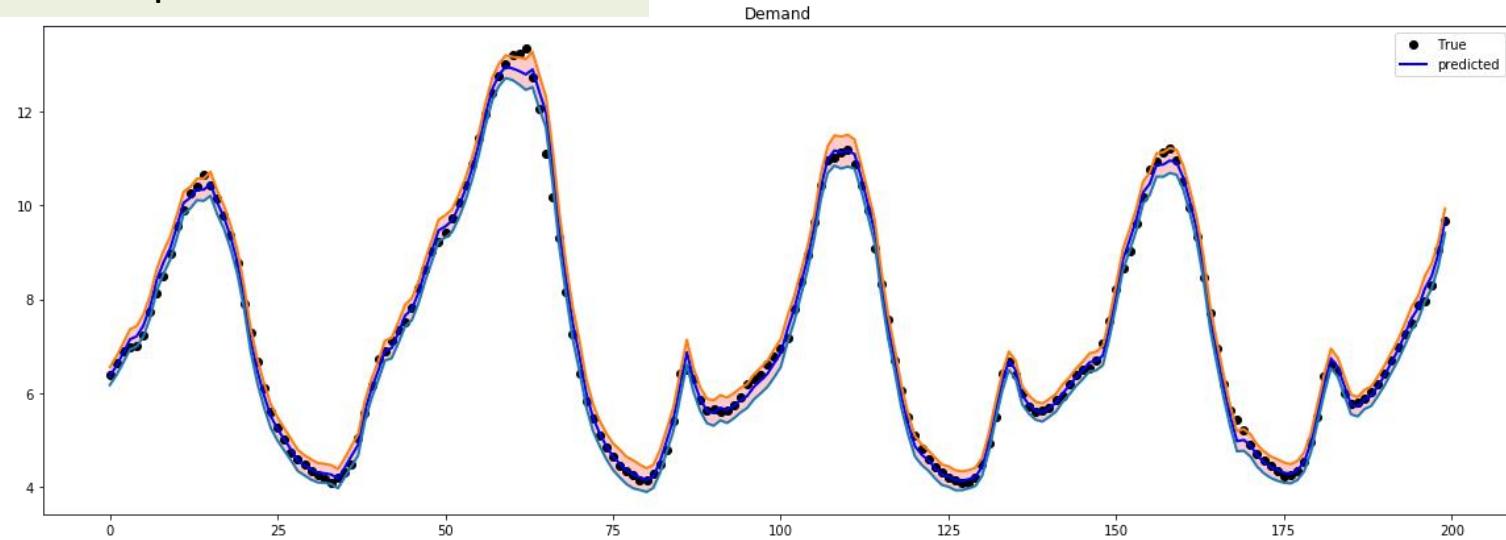
$$\hat{y}(t+1) = g(\mathbf{x}(t+1), \mathbf{u}(t+1)) + \mathbf{v}(t+1)$$

Systems Identification

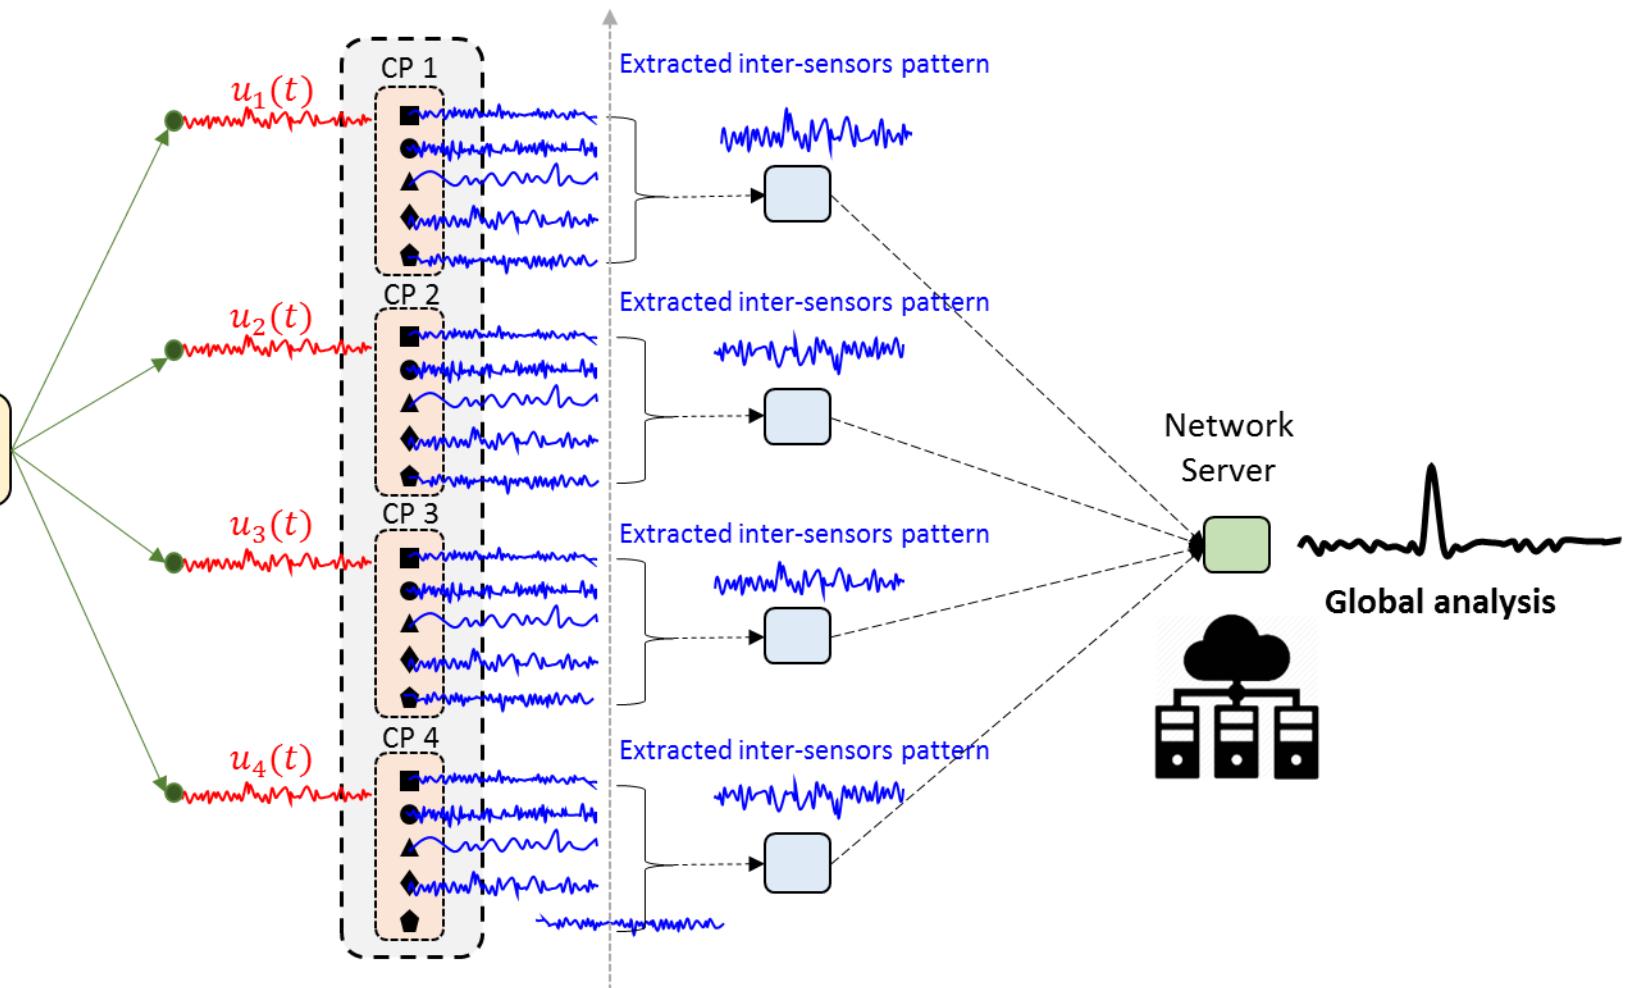
Wind farm power generation prediction



Energy demand prediction



Systems Monitoring



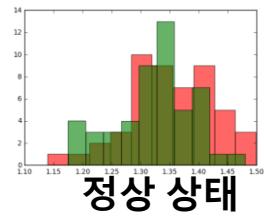
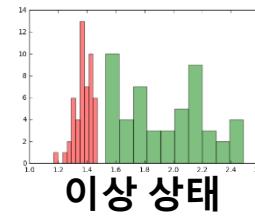
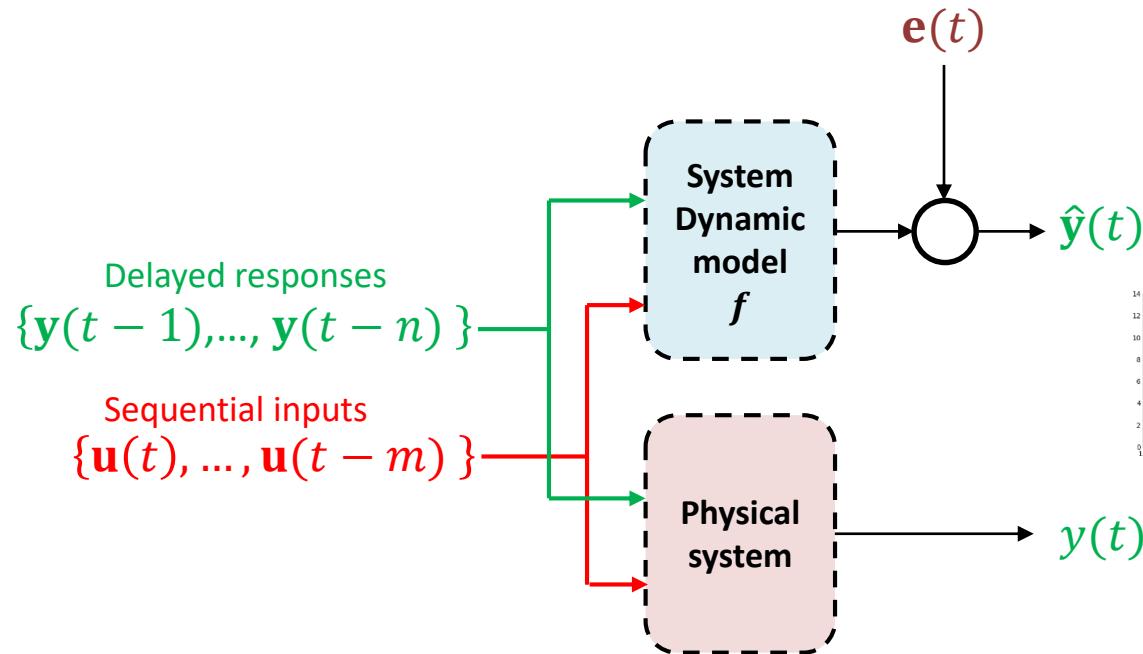
IoT Hierarchical
Data processing :



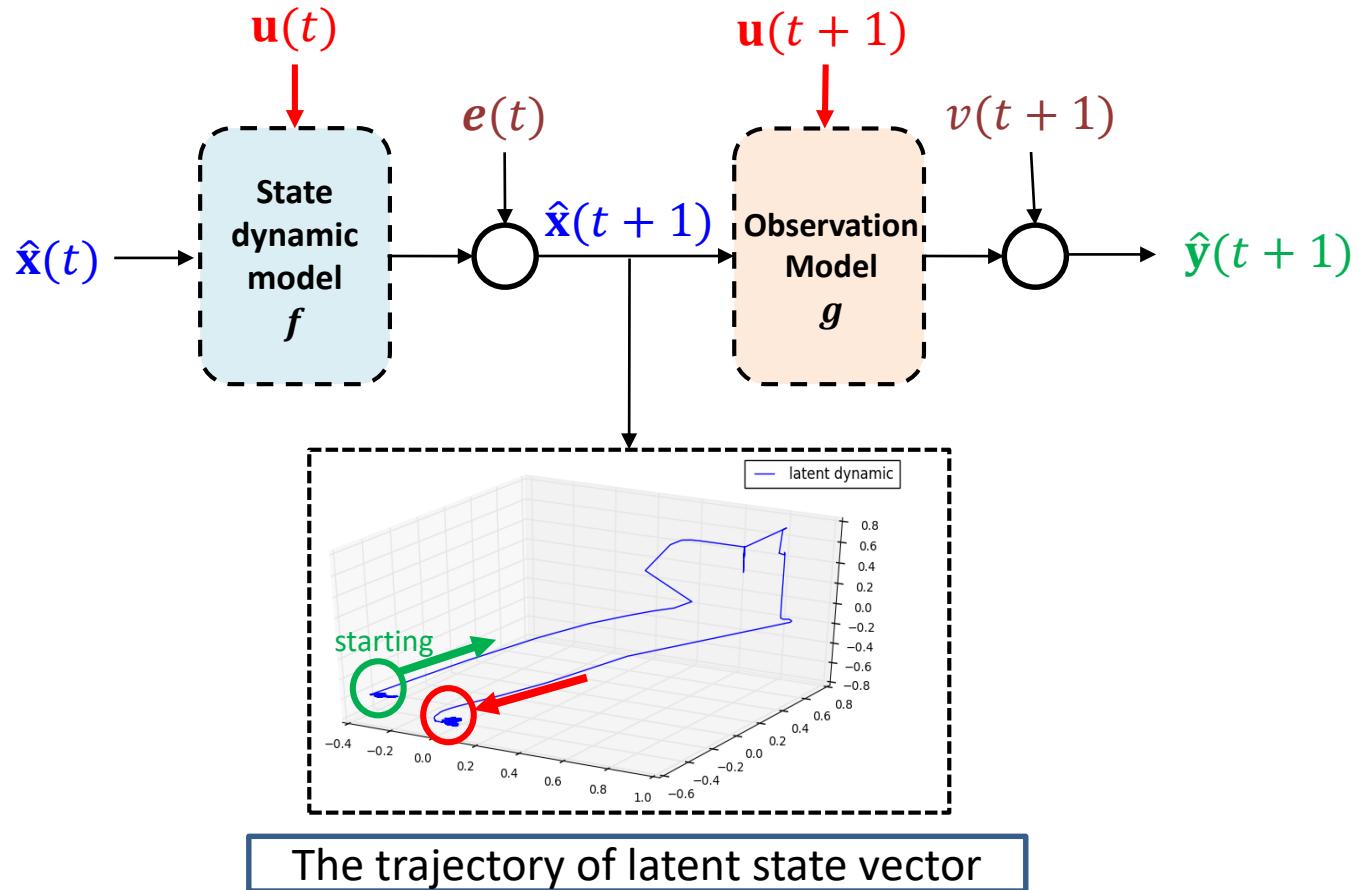
AI Hierarchical
machine intelligence :



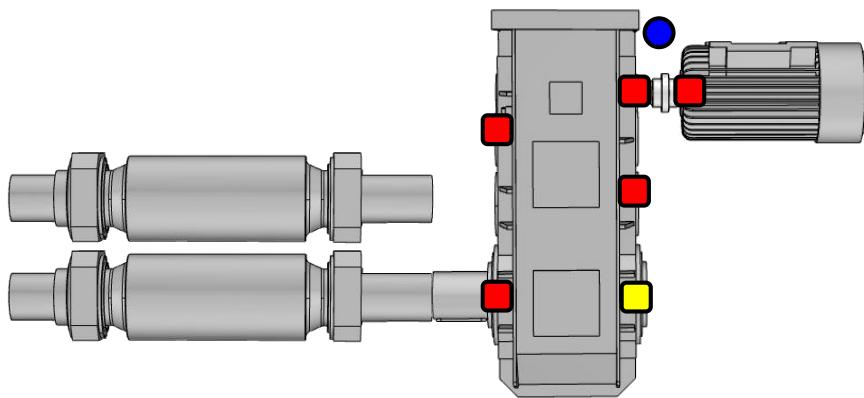
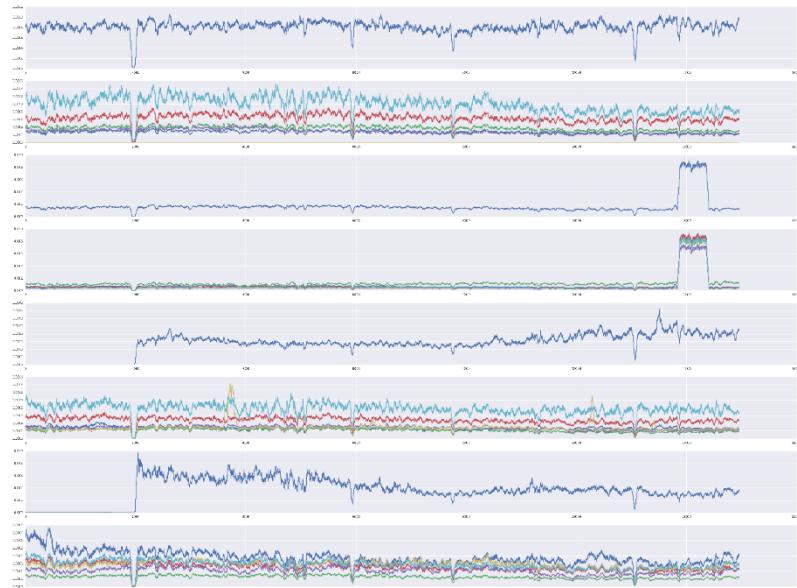
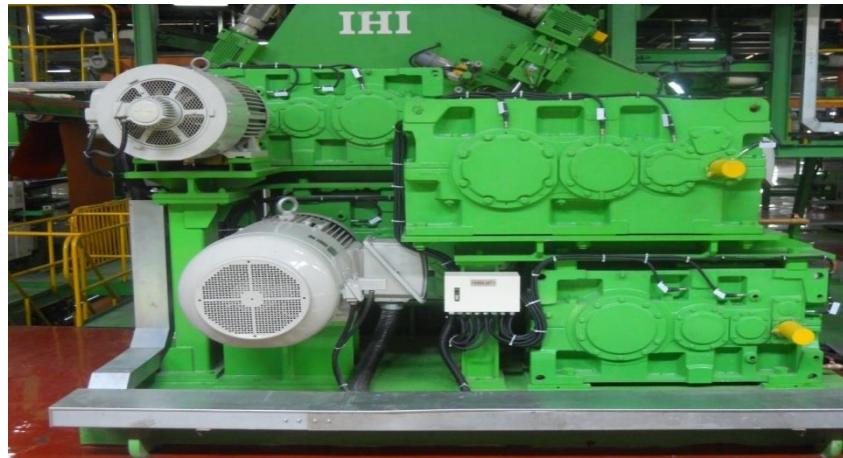
Systems Monitoring



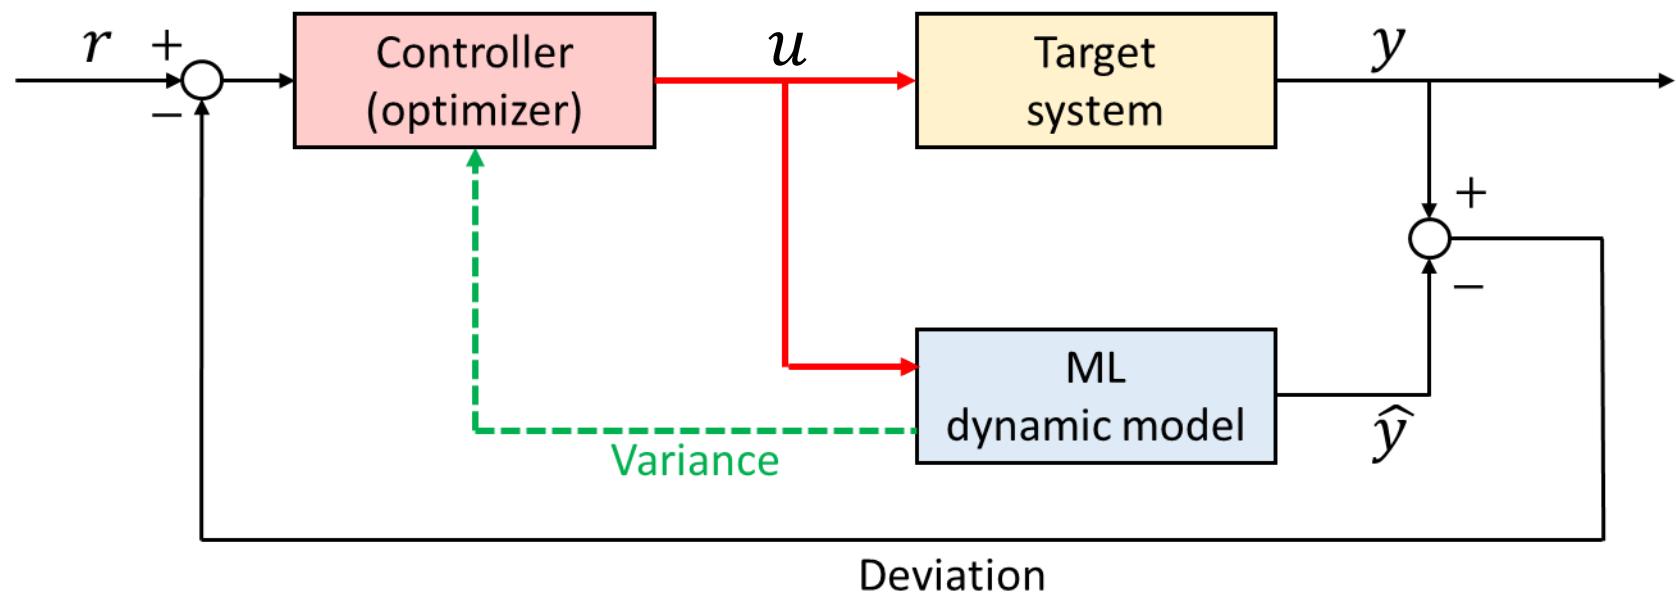
Systems Monitoring

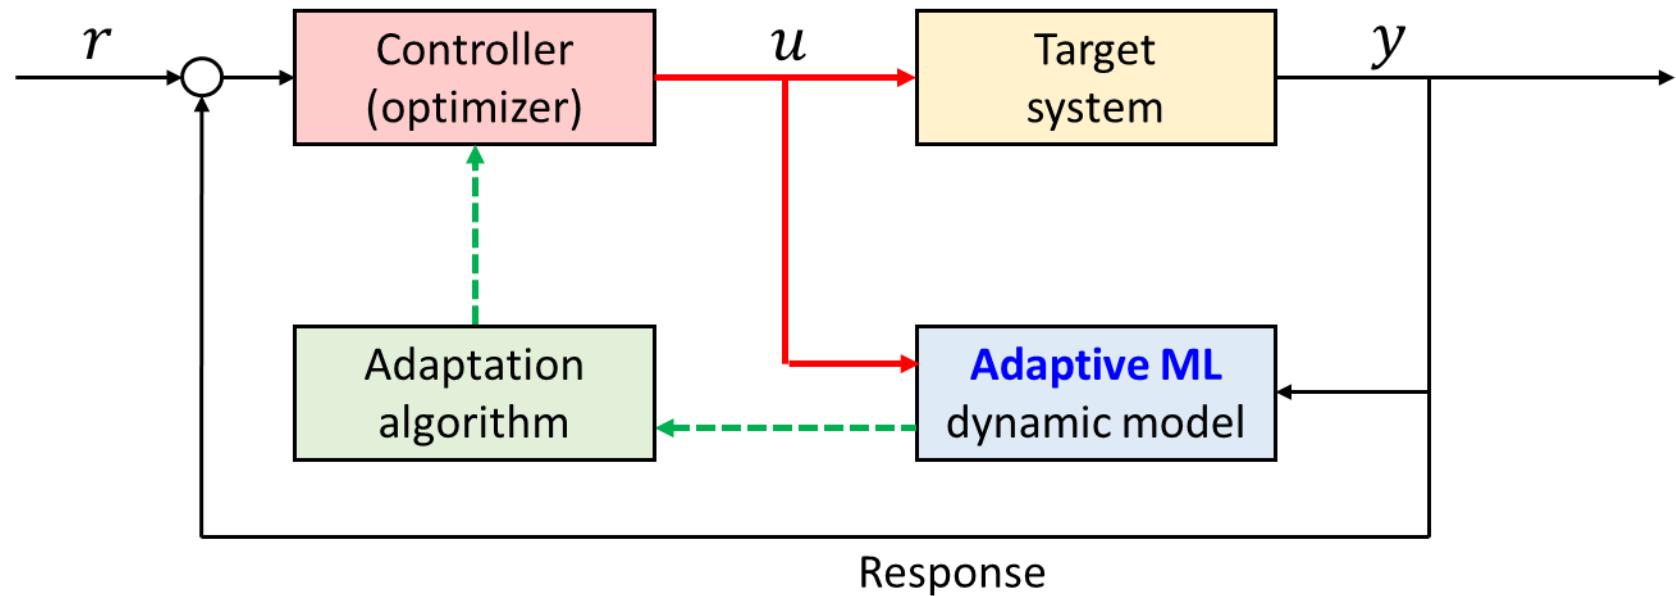


Systems Monitoring

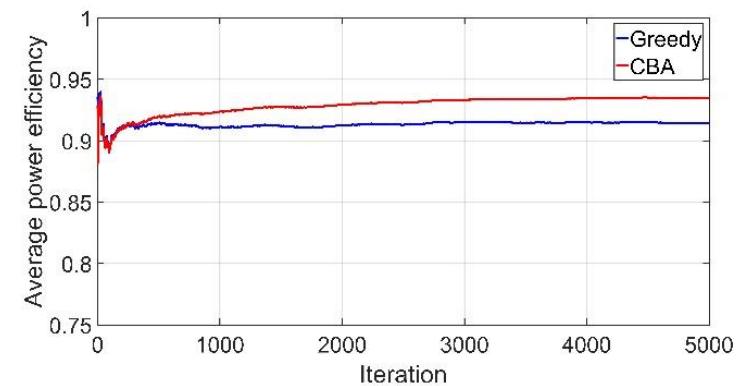
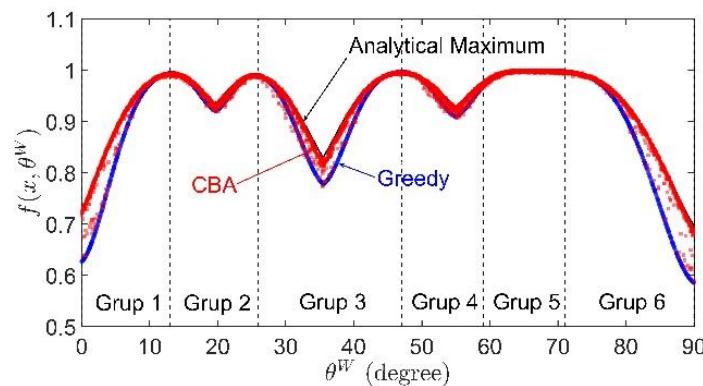
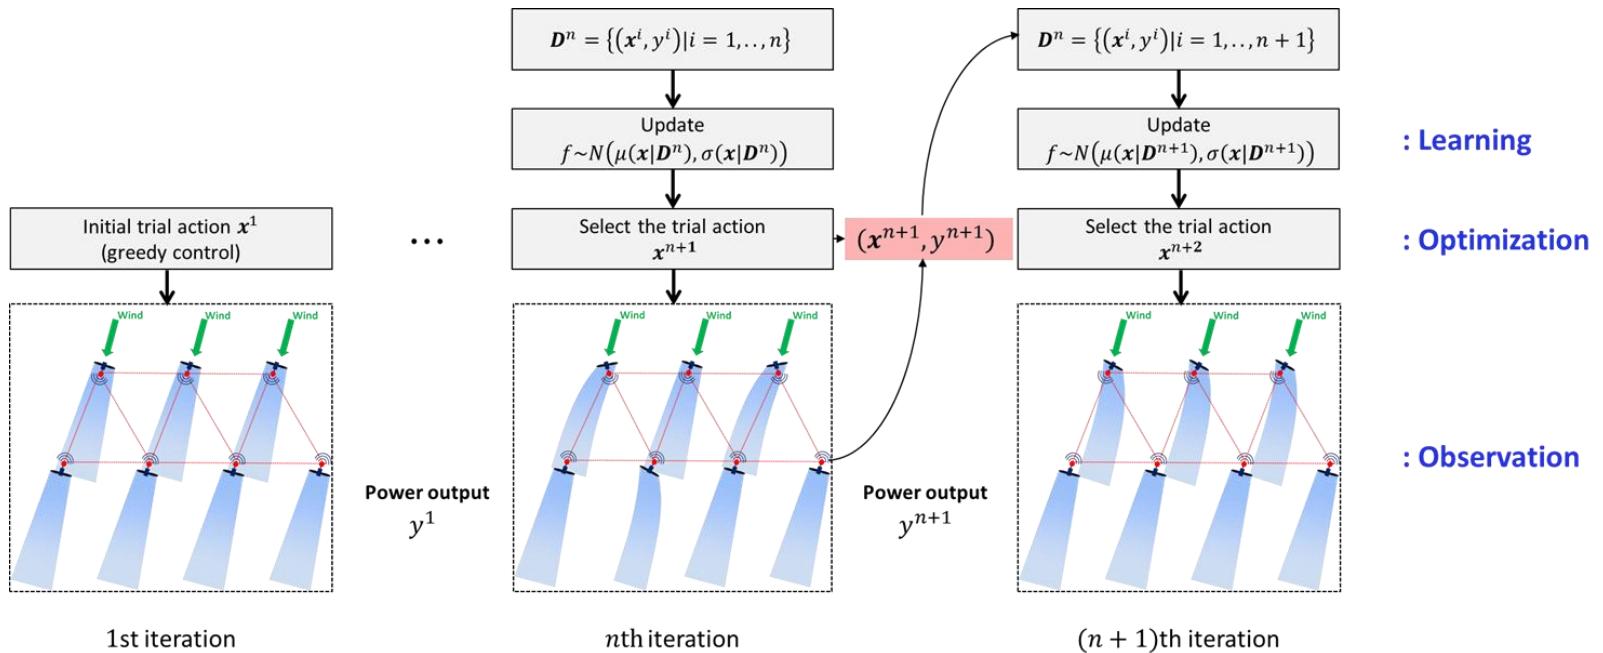


Systems Control

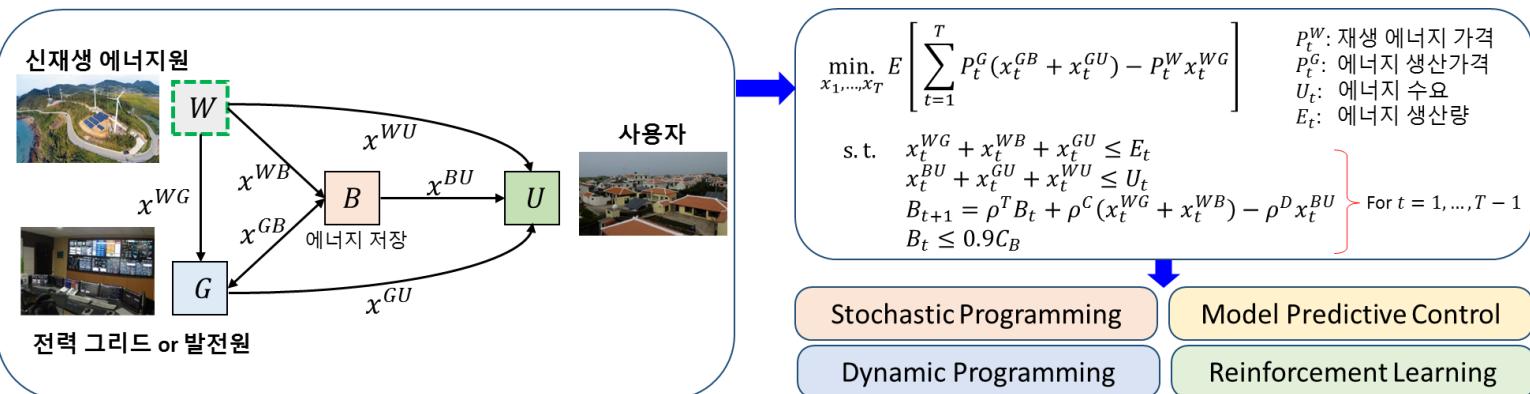
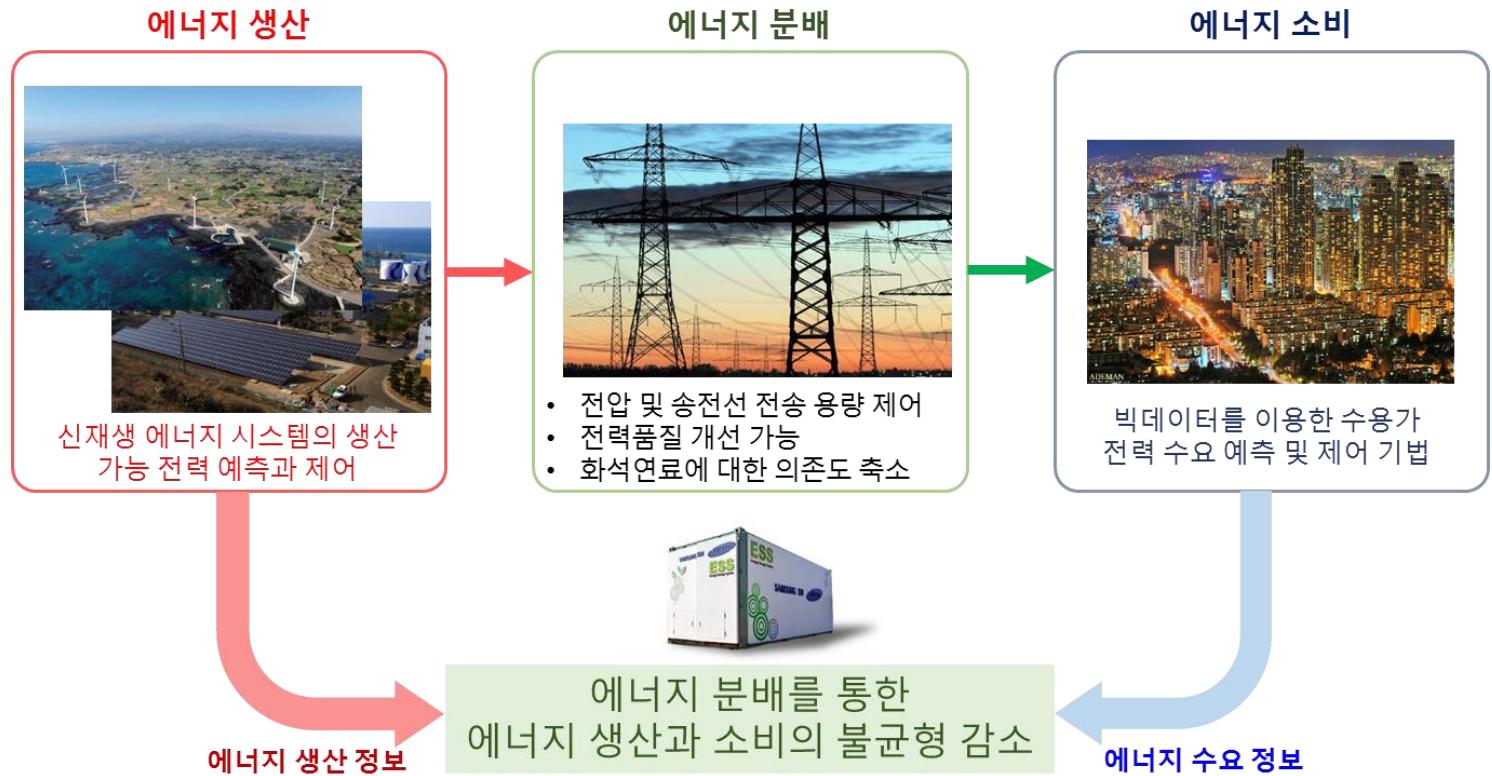




Systems Control



Systems Control



Course schedule

1. Bayesian Modeling and Inference (3 weeks)

- Probability distributions
- Prior, Likelihood, and Posterior
- Conjugate models
- Hierarchical Models
- Elements of Computational Bayesian Statistics

2. Single-agent, single-stage decision-making (3 weeks)

- Bayesian regression
- Bayesian classification
- Bayesian Network
- Influential Diagram

3. Multi-stages decision-making (6 weeks)

- Bandit problem
- Bayesian Optimization
- Markov Decision Process (MDP)
- Dynamic Programming
- Reinforcement Learning
 - Monte Carlo Methods
 - Temporal Difference Methods

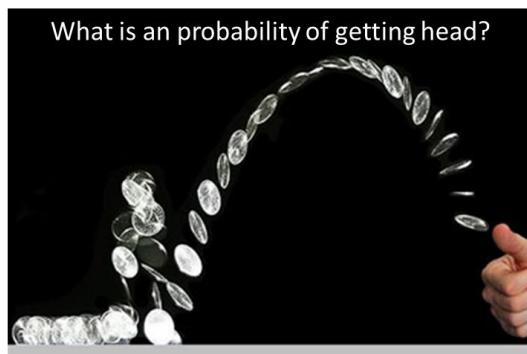
4. Multi-agents decision-making (3 weeks)

- Basics of game theory
- Bayesian Game (with uncertainty about other agents)

Main subject

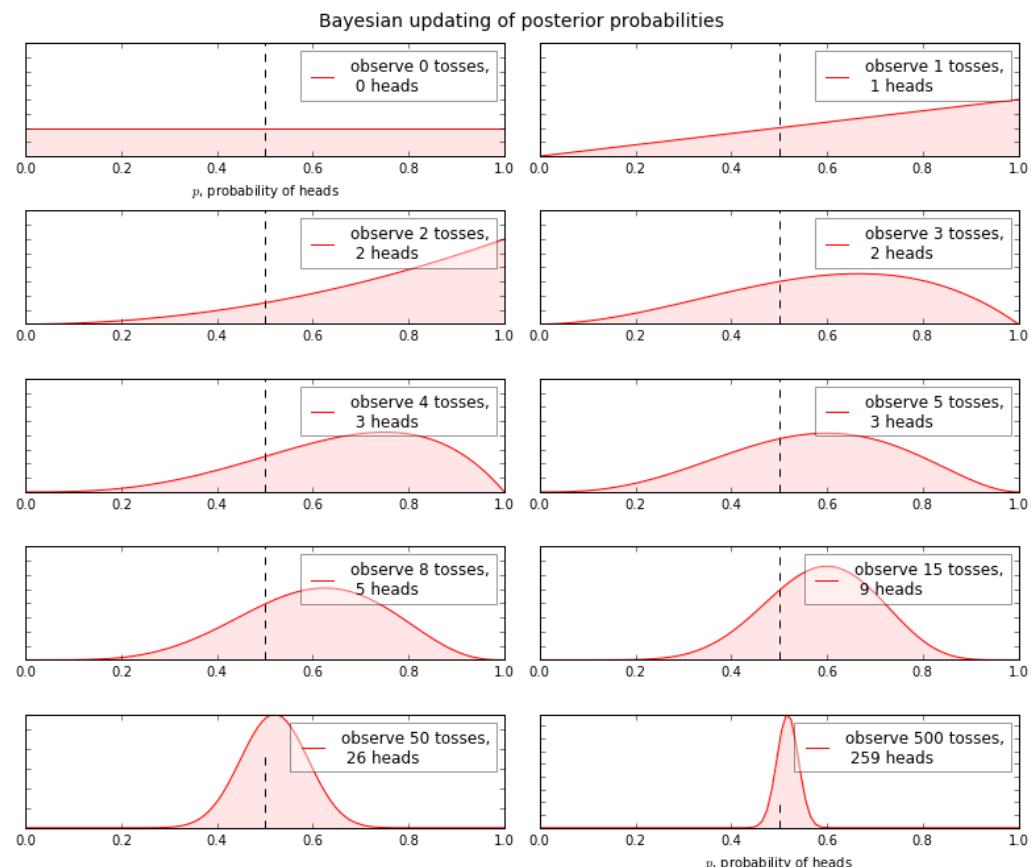
1. Bayesian Modeling and Inference (3 weeks)

- Probability distributions
- **Prior, Likelihood, and Posterior**
- Conjugate models
- Hierarchical Models
- Elements of Computational Bayesian Statistics



$$\begin{aligned} p(\theta|y) &= \frac{p(y|\theta)p(\theta)}{p(y)} \\ &= \frac{p(y|\theta)p(\theta)}{\int_{\theta} p(y|\theta)p(\theta)d\theta} \\ &\propto p(y|\theta)p(\theta) \end{aligned}$$

Posterior \propto Likelihood X Prior

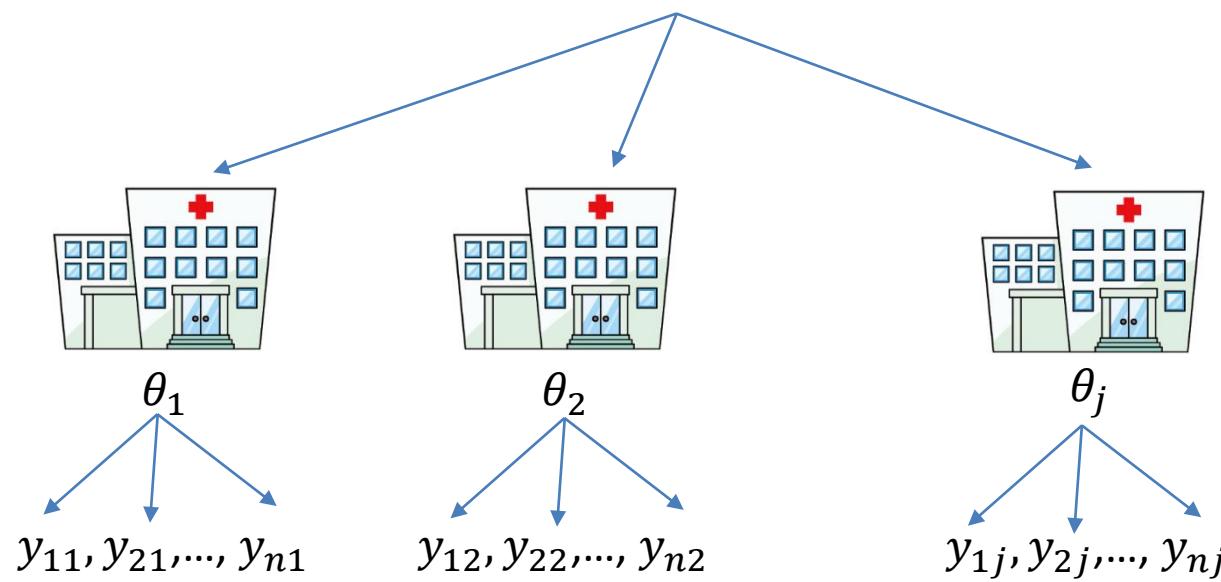


Main subject

1. Bayesian Modeling and Inference (3 weeks)

- Probability distributions
- Prior, Likelihood, and Posterior
- Conjugate models
- **Hierarchical Models**

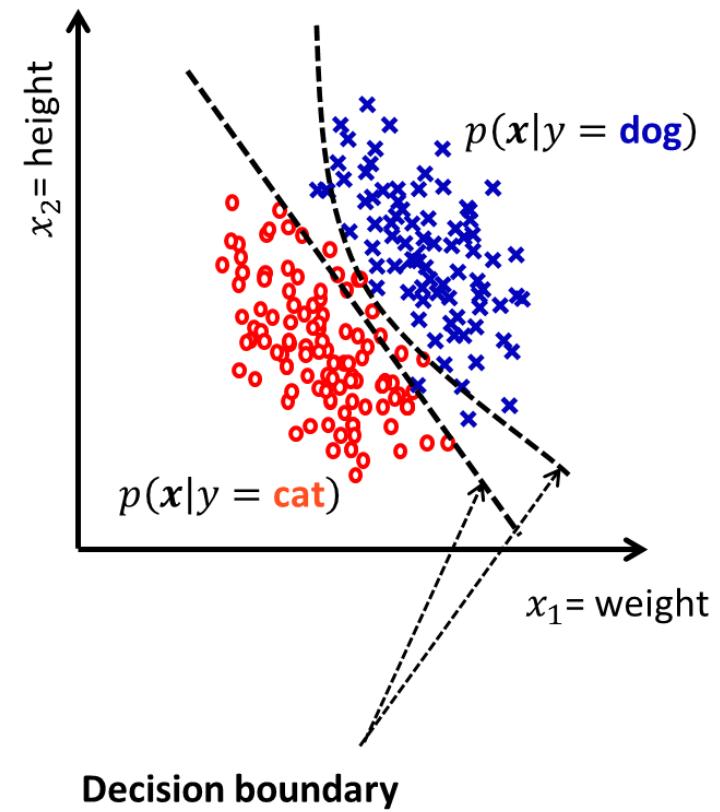
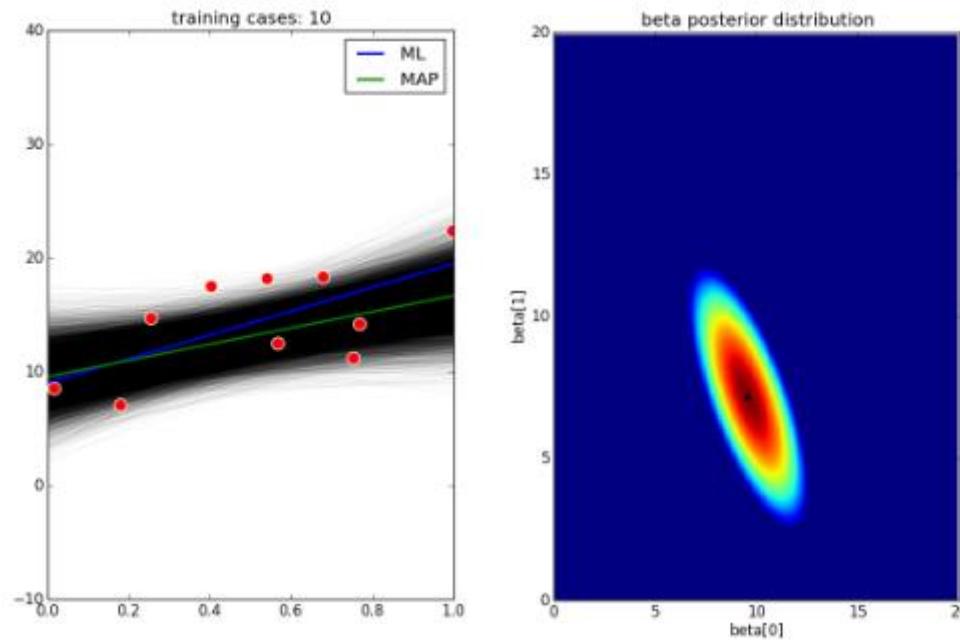
Survival probability of cardiac patients $\theta_j \sim$ population distribution



Main subject

2. Single-agent, single-stage decision-making (3 weeks)

- Bayesian regression
- Bayesian classification
- Bayesian Network
- Influential Diagram

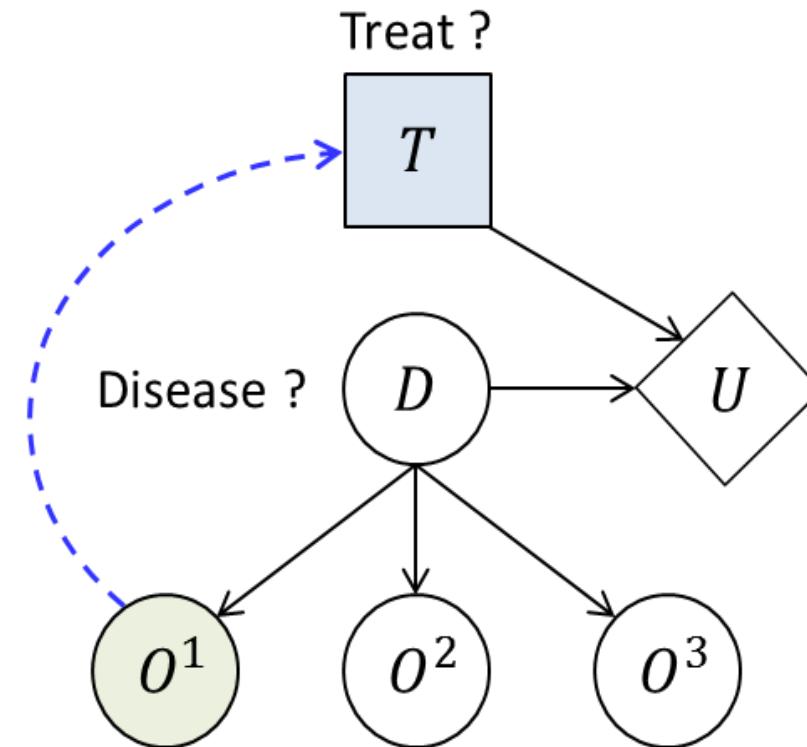


$$P(y = \text{dog}|x) = P(y = \text{cat}|x)$$

Main subject

2. Single-agent, single-stage decision-making (3 weeks)

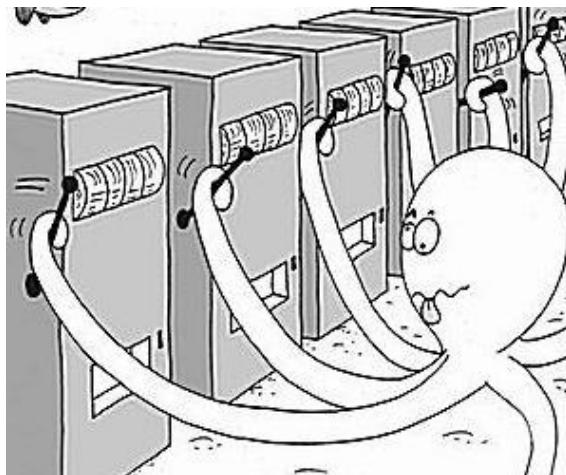
- Bayesian regression
- Bayesian classification
- **Bayesian Network**
- **Influential Diagram**



Main subject

3. Multi-stages decision-making (6 weeks)

- **Bandit problem**
- Bayesian Optimization
- Markov Decision Process (MDP)
- Dynamic Programming
- Reinforcement Learning
 - Monte Carlo Methods
 - Temporal Difference Methods



Facing with N slot machines with different payoff distribution,
→ **Devise a strategy to maximize payoff**
→ **Find the best bandit as quickly as possible**

Acquiring new information
(exploration)

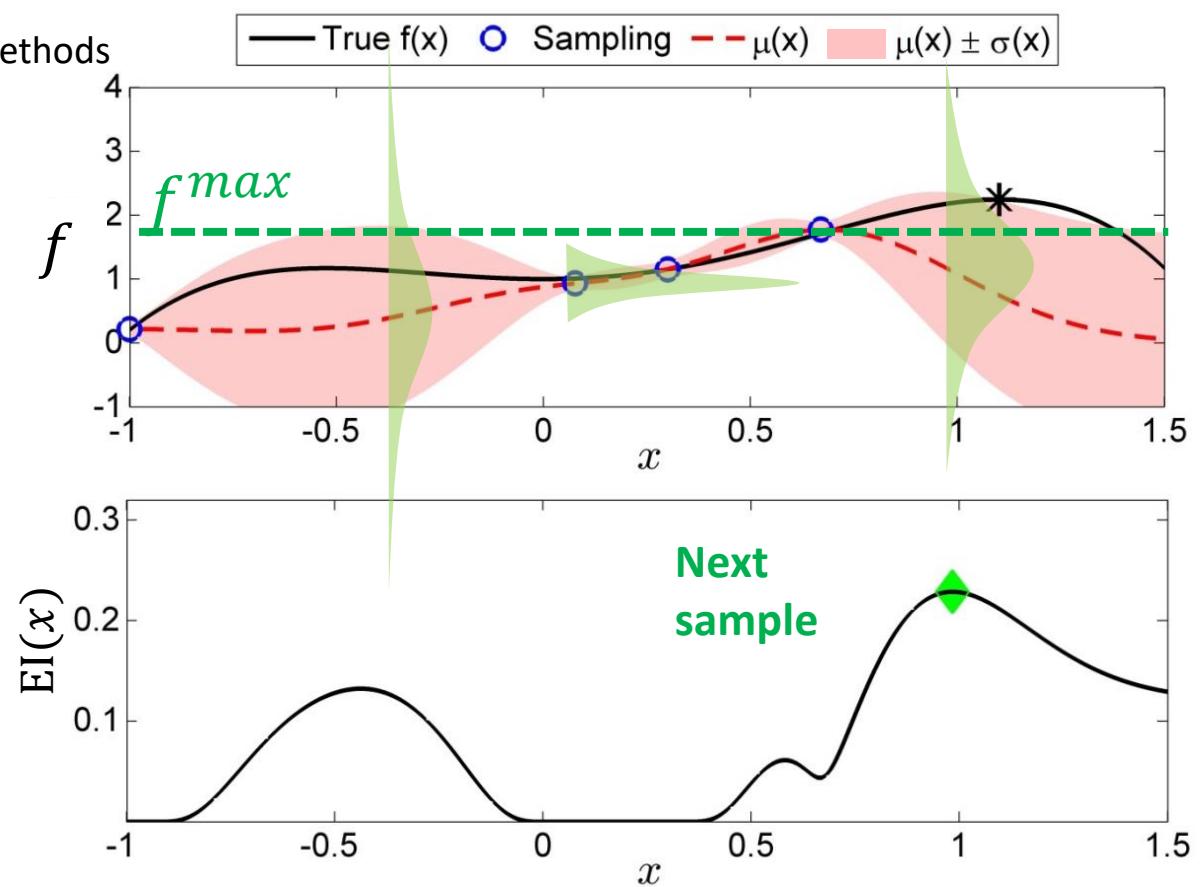
trade-off

capitalizing on the information
available so far
(exploitation)

Main subject

3. Multi-stages decision-making (6 weeks)

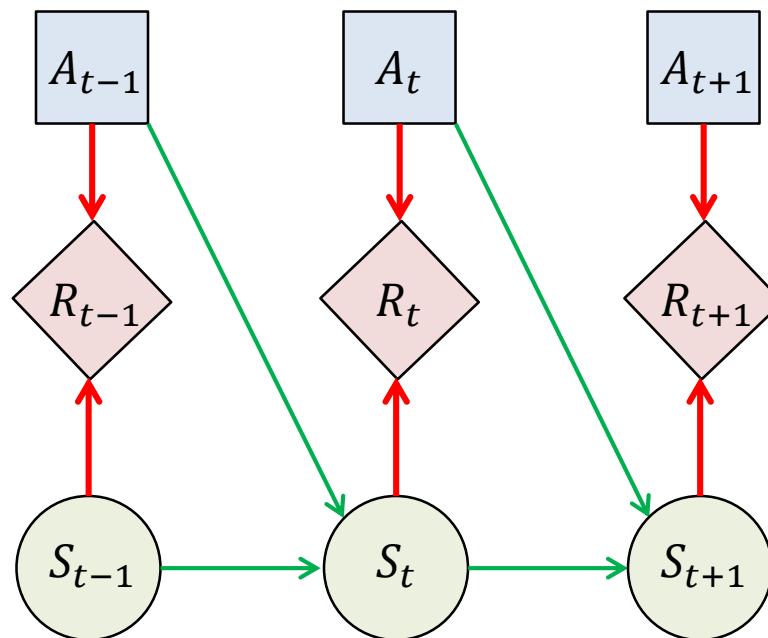
- Bandit problem
- **Bayesian Optimization**
- Markov Decision Process (MDP)
- Dynamic Programming
- Reinforcement Learning
 - Monte Carlo Methods
 - Temporal Difference Methods



Main subject

3. Multi-stages decision-making (6 weeks)

- Bandit problem
- Bayesian Optimization
- **Markov Decision Process (MDP)**
- Dynamic Programming
- Reinforcement Learning
 - Monte Carlo Methods
 - Temporal Difference Methods



A_t : action taken at time t

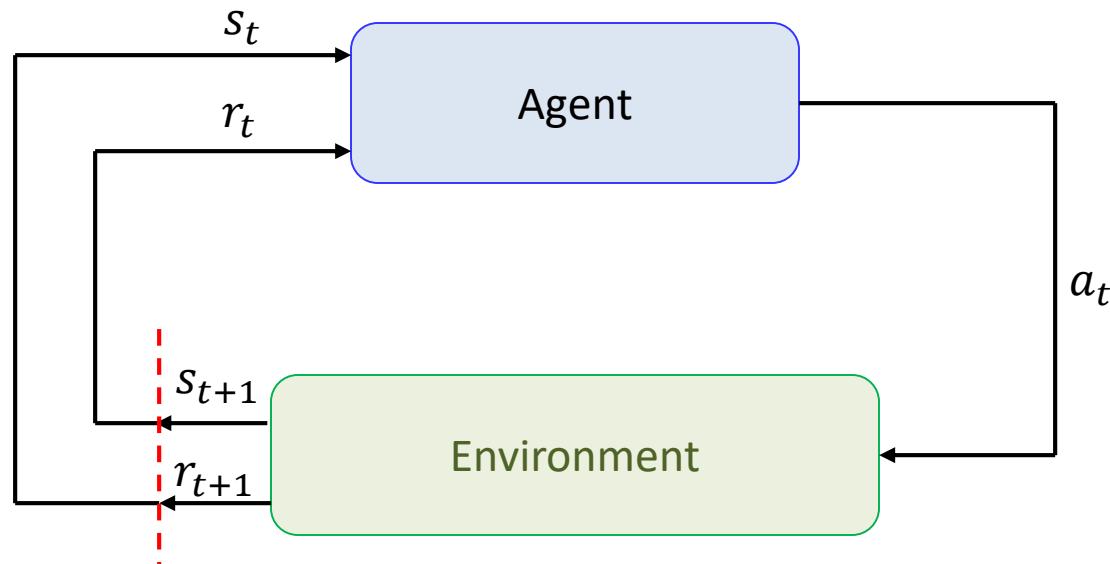
R_t : reward received at time t

S_t : state at time t

Main subject

3. Multi-stages decision-making (6 weeks)

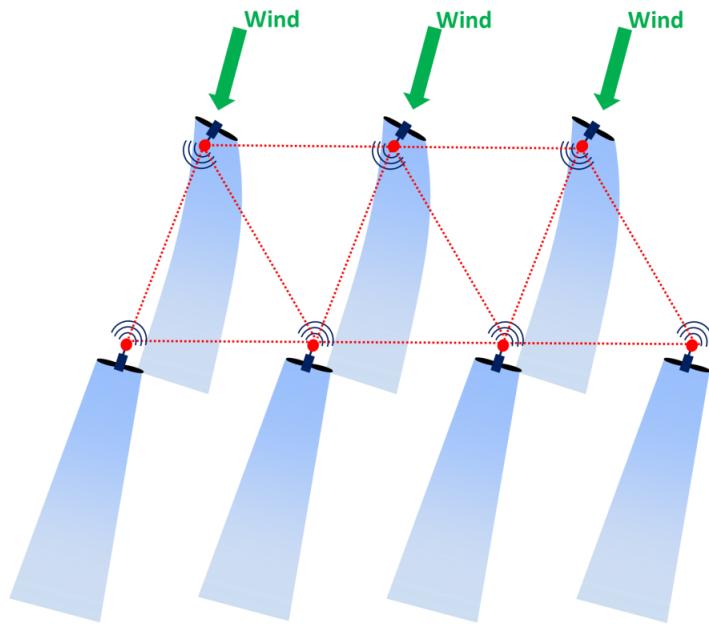
- Bandit problem
- Bayesian Optimization
- Markov Decision Process (MDP)
- Dynamic Programming
- **Reinforcement Learning**
 - Monte Carlo Methods
 - Temporal Difference Methods



Main subject

4. Multi-agents decision-making (3 weeks)

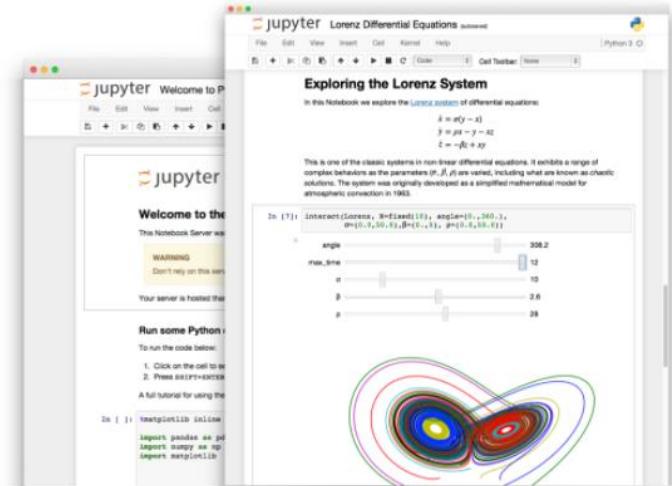
- **Basics of game theory**
- Bayesian Game (with uncertainty about other agents)



WT 1
greedy
 (α_1^G, o_1^G)
cooper
 (α_1^C, o_1^C)

WT 2	greedy (α_2^G, o_2^G)	cooper (α_2^C, o_2^C)
greedy (α_1^G, o_1^G)	$(70, 70)$ $\sum = 140$	$(70, 50)$ $\sum = 120$
cooper (α_1^C, o_1^C)	$(50, 70)$ $\sum = 120$	$(60, 90)$ $\sum = 150$

Computation tool



Language of choice

The Notebook has support for over 40 programming languages, including those popular in Data Science such as Python, R, Julia and Scala.



Share notebooks

Notebooks can be shared with others using email, Dropbox, GitHub and the Jupyter Notebook Viewer.



The Jupyter Notebook

The Jupyter Notebook is an open-source web application that allows you to create and share documents that contain live code, equations, visualizations and explanatory text. Uses include: data cleaning and transformation, numerical simulation, statistical modeling, machine learning and much more.

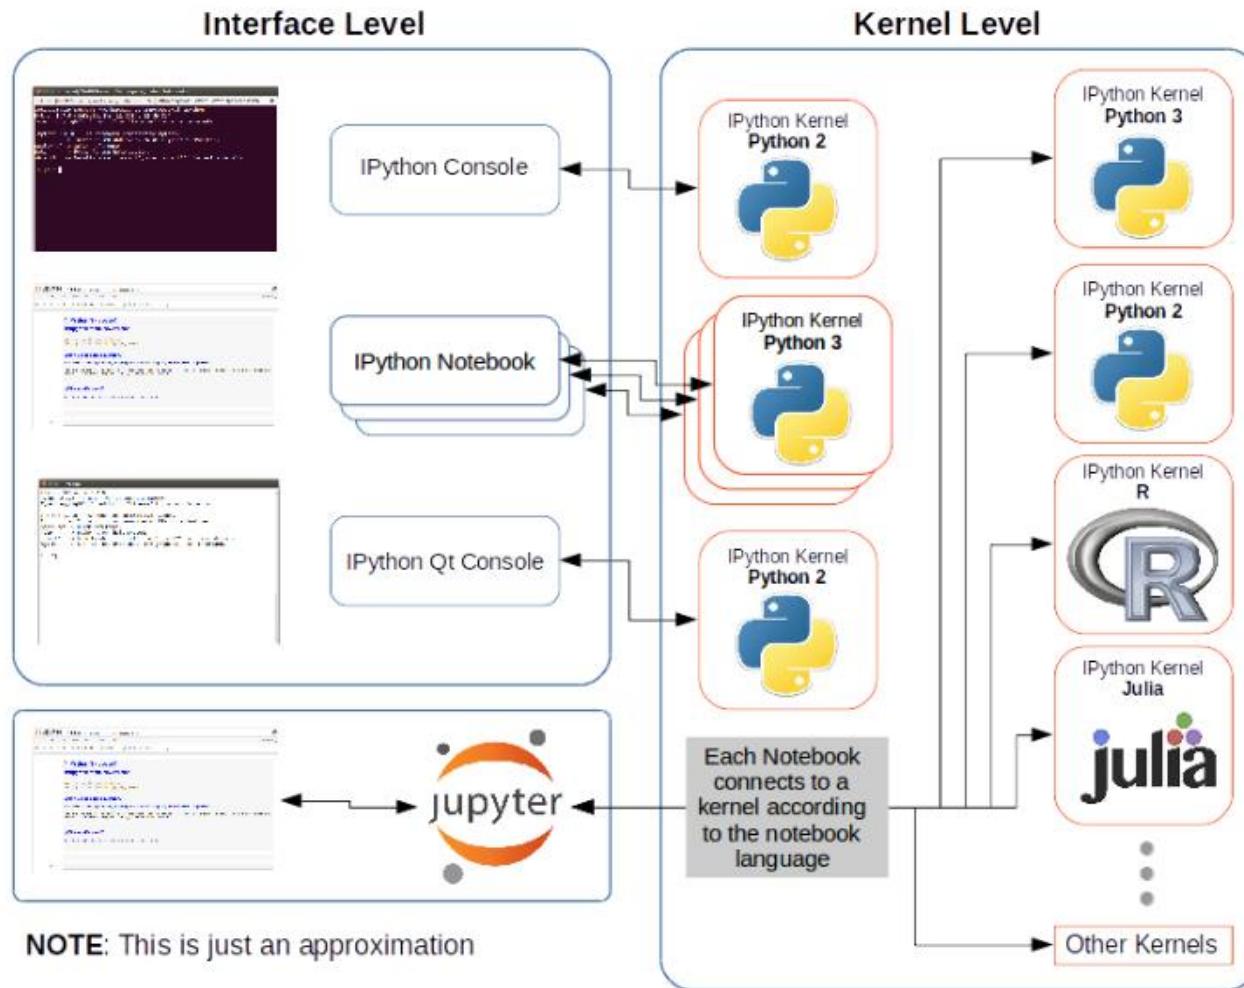


Big data integration

Leverage big data tools, such as Apache Spark, from Python, R and Scala. Explore that same data with pandas, scikit-learn, ggplot2, dplyr, etc.

- It combines compiler, ide, text editor, graphic editor...

Computation tool



- Jupyter is the latest version of Ipython and now it includes more than 10 kernels (i.e., python, R, Julia,...)

Computation tool

Basic scientific computing

modulus



NumPy

Base N-dimensional array package



SciPy library

Fundamental library for scientific computing



Matplotlib

Comprehensive 2D Plotting



IPython

Enhanced Interactive Console



Sympy

Symbolic mathematics



pandas

Data structures & analysis

Application oriented modules



[build](#) [travis](#) [coverage](#) 87% [powered by](#) NumFOCUS

PyMC3 is a Python package for Bayesian statistical modeling and Probabilistic Machine Learning which focuses on advanced Markov chain Monte Carlo and variational fitting algorithms. Its flexibility and extensibility make it applicable to a large suite of problems.

Edward



A library for probabilistic modeling, inference, and criticism.

Edward is a Python library for probabilistic modeling, inference, and criticism. It is a testbed for fast experimentation and research with probabilistic models, ranging from classical hierarchical models on small data sets to complex deep probabilistic models on large data sets. Edward fuses three fields: Bayesian statistics and machine learning, deep learning, and probabilistic programming.

It supports **modeling** with

GPy

The Gaussian processes framework in Python.

Computation tool



JUPYTER FAQ </> ⋮ ⌂ ⏪

In Depth: Linear Regression

Just as naive Bayes (discussed earlier in [In Depth: Naive Bayes Classification](#)) is a good starting point for classification tasks, linear regression models are a good starting point for regression tasks. Such models are popular because they can be fit very quickly, and are very interpretable. You are probably familiar with the simplest form of a linear regression model (i.e., fitting a straight line to data) but such models can be extended to model more complicated data behavior.

In this section we will start with a quick intuitive walk-through of the mathematics behind this well-known problem, before seeing how linear models can be generalized to account for more complicated patterns in data.

We begin with the standard imports:

Text regions:

Import Python modulus:

Code for plotting:

Result plot

```
In [1]: %matplotlib inline  
import matplotlib.pyplot as plt  
import seaborn as sns; sns.set()  
import numpy as np
```

Simple Linear Regression

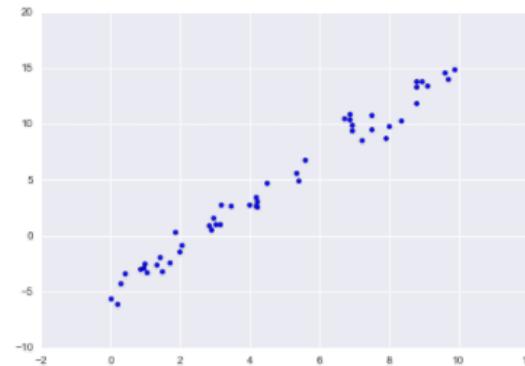
We will start with the most familiar linear regression, a straight-line fit to data. A straight-line fit is a model of the form

$$y = ax + b$$

where a is commonly known as the *slope*, and b is commonly known as the *intercept*.

Consider the following data, which is scattered about a line with a slope of 2 and an intercept of -5:

```
In [2]: rng = np.random.RandomState(1)  
x = 10 * rng.rand(50)  
y = 2 * x - 5 + rng.randn(50)  
plt.scatter(x, y);
```



We can use Scikit-Learn's `LinearRegression` estimator to fit this data and construct the best-fit line:

```
In [3]: from sklearn.linear_model import LinearRegression  
model = LinearRegression(fit_intercept=True)  
model.fit(x[:, np.newaxis], y)
```

Computation tool

Prerequisite: Python

While Jupyter runs code in many programming languages, **Python** is a requirement (Python 3.3 or greater, or Python 2.7) for installing the Jupyter Notebook.

We recommend using the [Anaconda](#) distribution to install Python and Jupyter. We'll go through its installation in the next section.

Installing Jupyter using Anaconda and conda

For new users, we **highly recommend** installing Anaconda. Anaconda conveniently installs Python, the Jupyter Notebook, and other commonly used packages for scientific computing and data science.

We recommend using the Anaconda distribution to install Python and Jupyter. We'll go through its installation in the next section.

- Download [Anaconda](#). We recommend downloading Anaconda's latest Python 3 version (currently Python 3.5).
- Install the version of Anaconda which you downloaded, following the instructions on the download page.
- Congratulations, you have installed Jupyter Notebook. To run the notebook:

```
jupyter notebook
```

See [Running the Notebook](#) for more details.

Alternative for experienced Python users: Installing Jupyter with pip

As an existing Python user, you may wish to install Jupyter using Python's package manager, [pip](#), instead of Anaconda.

First, ensure that you have the latest pip; older versions may have trouble with some dependencies:

```
pip3 install --upgrade pip
```

Then install the Jupyter Notebook using:

```
pip3 install jupyter
```

(Use pip if using legacy Python 2.)

Congratulations. You have installed Jupyter Notebook.

See [Running the Notebook](#) for more details.

Computation tool

Anaconda 4.4.0 For Windows Installer

Python 3.6 version *

64-Bit Graphical Installer (437 MB) [?](#)

 Download

[Download 32-Bit \(362 MB\)](#)

Python 2.7 version *

64-Bit Graphical Installer (430 MB) [?](#)

 Download

[Download 32-Bit \(354 MB\)](#)

- If Anaconda is used, most scientific computing python modules (i.e., Scipy, Numpy, Matplotlib, Scikit, will be automatically installed)

Computation tool

```
Anaconda Prompt - jupyter notebook

(C:\Users\Jinkyoo\Anaconda2) C:\Users\Jinkyoo>cd JupyterTutorial
(C:\Users\Jinkyoo\Anaconda2) C:\Users\Jinkyoo\JupyterTutorial>dir
 Volume in drive C is Windows
 Volume Serial Number is 465D-6BE5

 Directory of C:\Users\Jinkyoo\JupyterTutorial

07/30/2017  02:24 AM    <DIR>      .
07/30/2017  02:24 AM    <DIR>      ..
07/30/2017  02:24 AM    <DIR>      .ipynb_checkpoints
09/16/2016  04:54 PM           425,506 L1_probabilityDistributions.ipynb
               1 File(s)     425,506 bytes
               3 Dir(s)   95,098,327,040 bytes free

(C:\Users\Jinkyoo\Anaconda2) C:\Users\Jinkyoo\JupyterTutorial>jupyter notebook
[I 02:26:47.309 NotebookApp] Serving notebooks from local directory: C:\Users\Jinkyoo\JupyterTutorial
[I 02:26:47.311 NotebookApp] 0 active kernels
[I 02:26:47.312 NotebookApp] The Jupyter Notebook is running at: http://localhost:8888/?token=8cd5dc5aa06b155aa83e7a209bf008abb96aa53dae7410d6
[I 02:26:47.315 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[C 02:26:47.355 NotebookApp]

Copy/paste this URL into your browser when you connect for the first time,
to login with a token:
 http://localhost:8888/?token=8cd5dc5aa06b155aa83e7a209bf008abb96aa53dae7410d6
[I 02:26:47.726 NotebookApp] Accepting one-time-token-authenticated connection from ::1
[W 02:26:47.951 NotebookApp] 404 GET /favicon.ico (::1) 11.00ms referer=None
```

Example: Data-driven wind farm power maximization

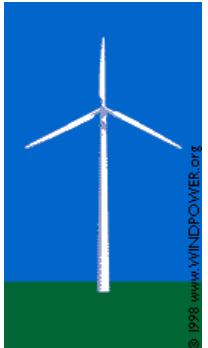
Size of wind farm ↑

Wind farm power production efficiency ↓

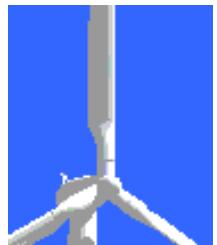
wind turbines	: companies	: assets
wind farm	: market	: stock market
wind:	resource	: money
Power:	revenue	: payoff

Photo courtesy of Vattenfall.

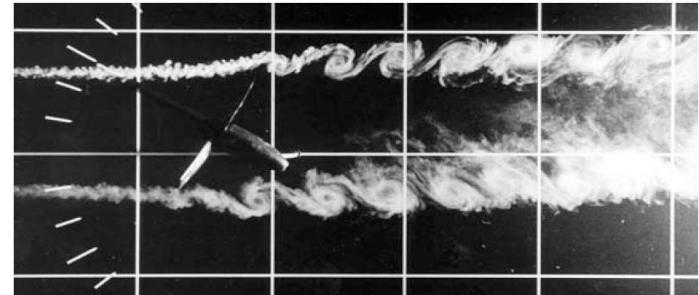
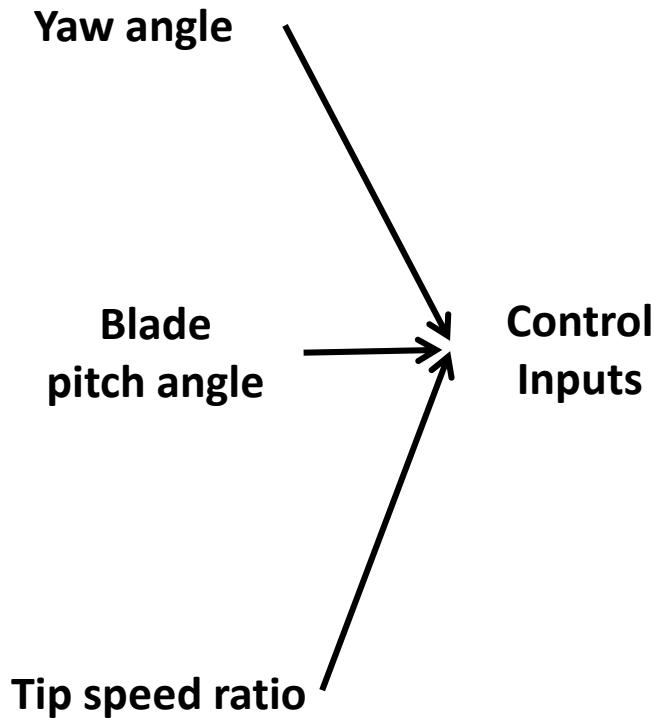
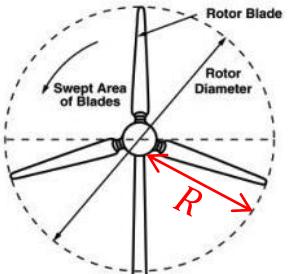
Example: Data-driven wind farm power maximization



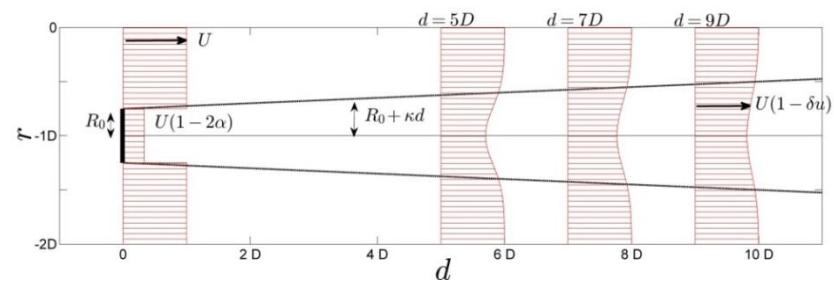
<www.windpowr.org>



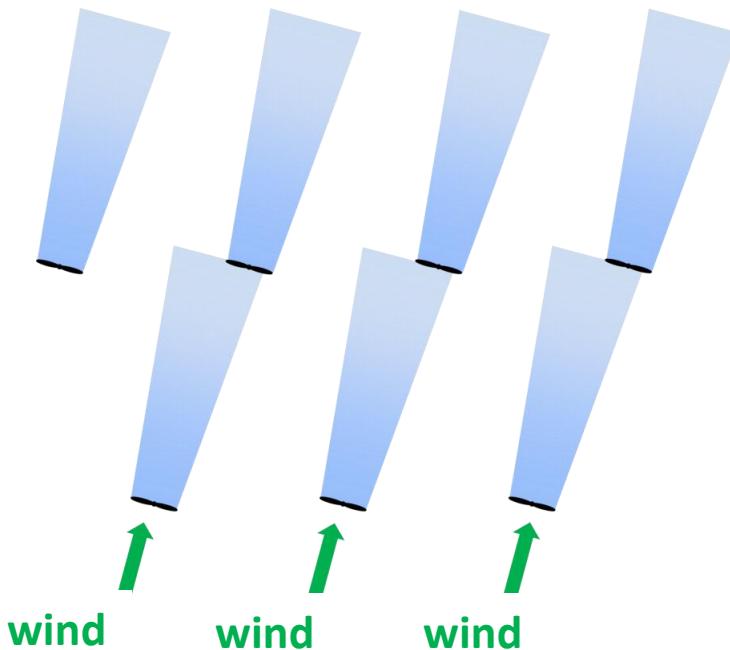
<www.windpowr.org>



Flow visualization of wake deflection. Conducted at the Royal Institute of Technology at 1987, Dahlberg (2003)



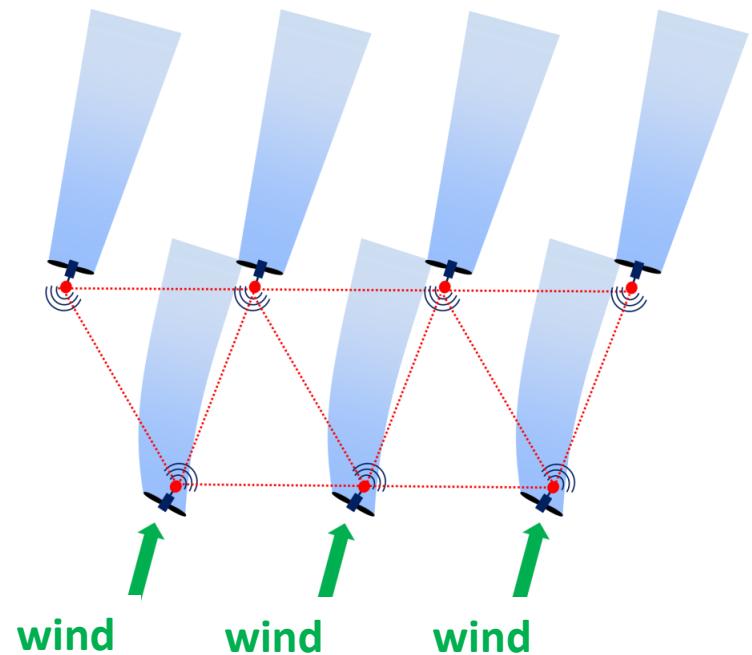
Example: Data-driven wind farm power maximization



Greedy strategy:

$$\sum_{i=1}^N \underset{x_i}{\text{maximize}} P_i(x; U, \theta)$$

•sensors



Cooperative strategy:

$$\underset{x}{\text{maximize}} \sum_{i=1}^N P_i(x; U, \theta)$$

$P_i(x; U, \theta)$:Power of wind turbine i

$x = (x_1, \dots, x_N)$: Control inputs for N wind turbines

Example: Data-driven wind farm power maximization

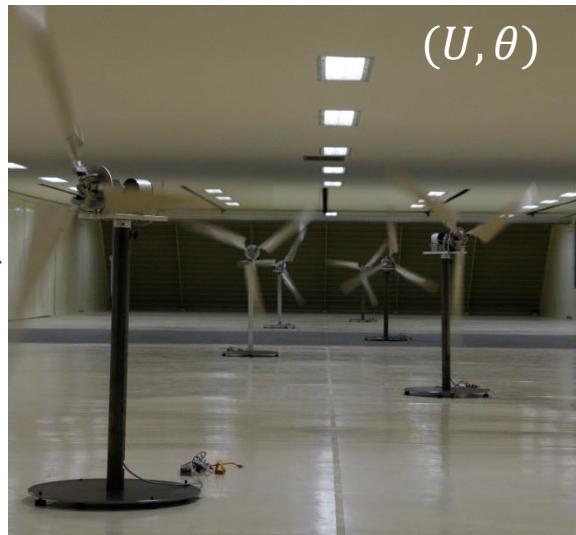
(U, θ) : Fixed wind condition (Fixed context)

Input (control actions)

$$\boldsymbol{x}^1 = (x_1^1, \dots, x_N^1)$$

$$\boldsymbol{x}^2 = (x_1^2, \dots, x_N^2)$$

\vdots



Output (power measurements)

$$y^1 = \sum_{i=1}^N P_i(\boldsymbol{x}^1; \theta, U) + \epsilon^1$$

$$y^2 = \sum_{i=1}^N P_i(\boldsymbol{x}^2; \theta, U) + \epsilon^2$$

\vdots

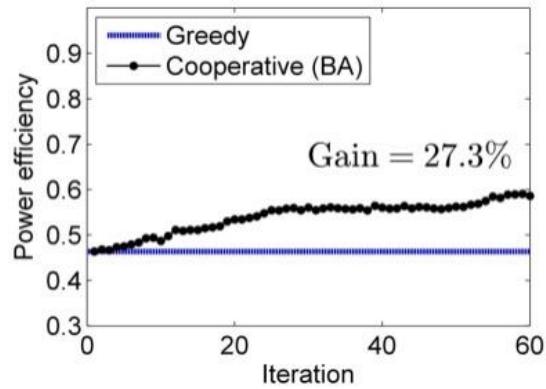
(y = total wind farm power)

subject to $\boldsymbol{x}^l \leq \boldsymbol{x} \leq \boldsymbol{x}^u$

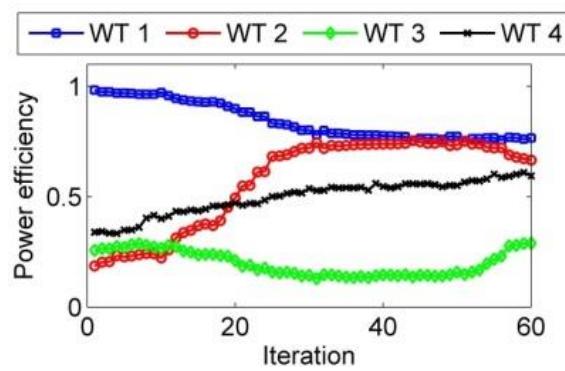
$$\underset{\boldsymbol{x}}{\text{maximize}} \ f(\boldsymbol{x}; \mathcal{U}, \theta) = \sum_{i=1}^N P_i(\boldsymbol{x}; \theta, U)$$

Without knowing $f(\boldsymbol{x}; \mathcal{U}, \theta)$

Example: Data-driven wind farm power maximization



Greedy control (initial)



Cooperative control (after convergence)

